

Security Best Practices for Amazon EKS

Hoseok Seo (he/him)
Solutions Architect



“Security is everyone’s job.”

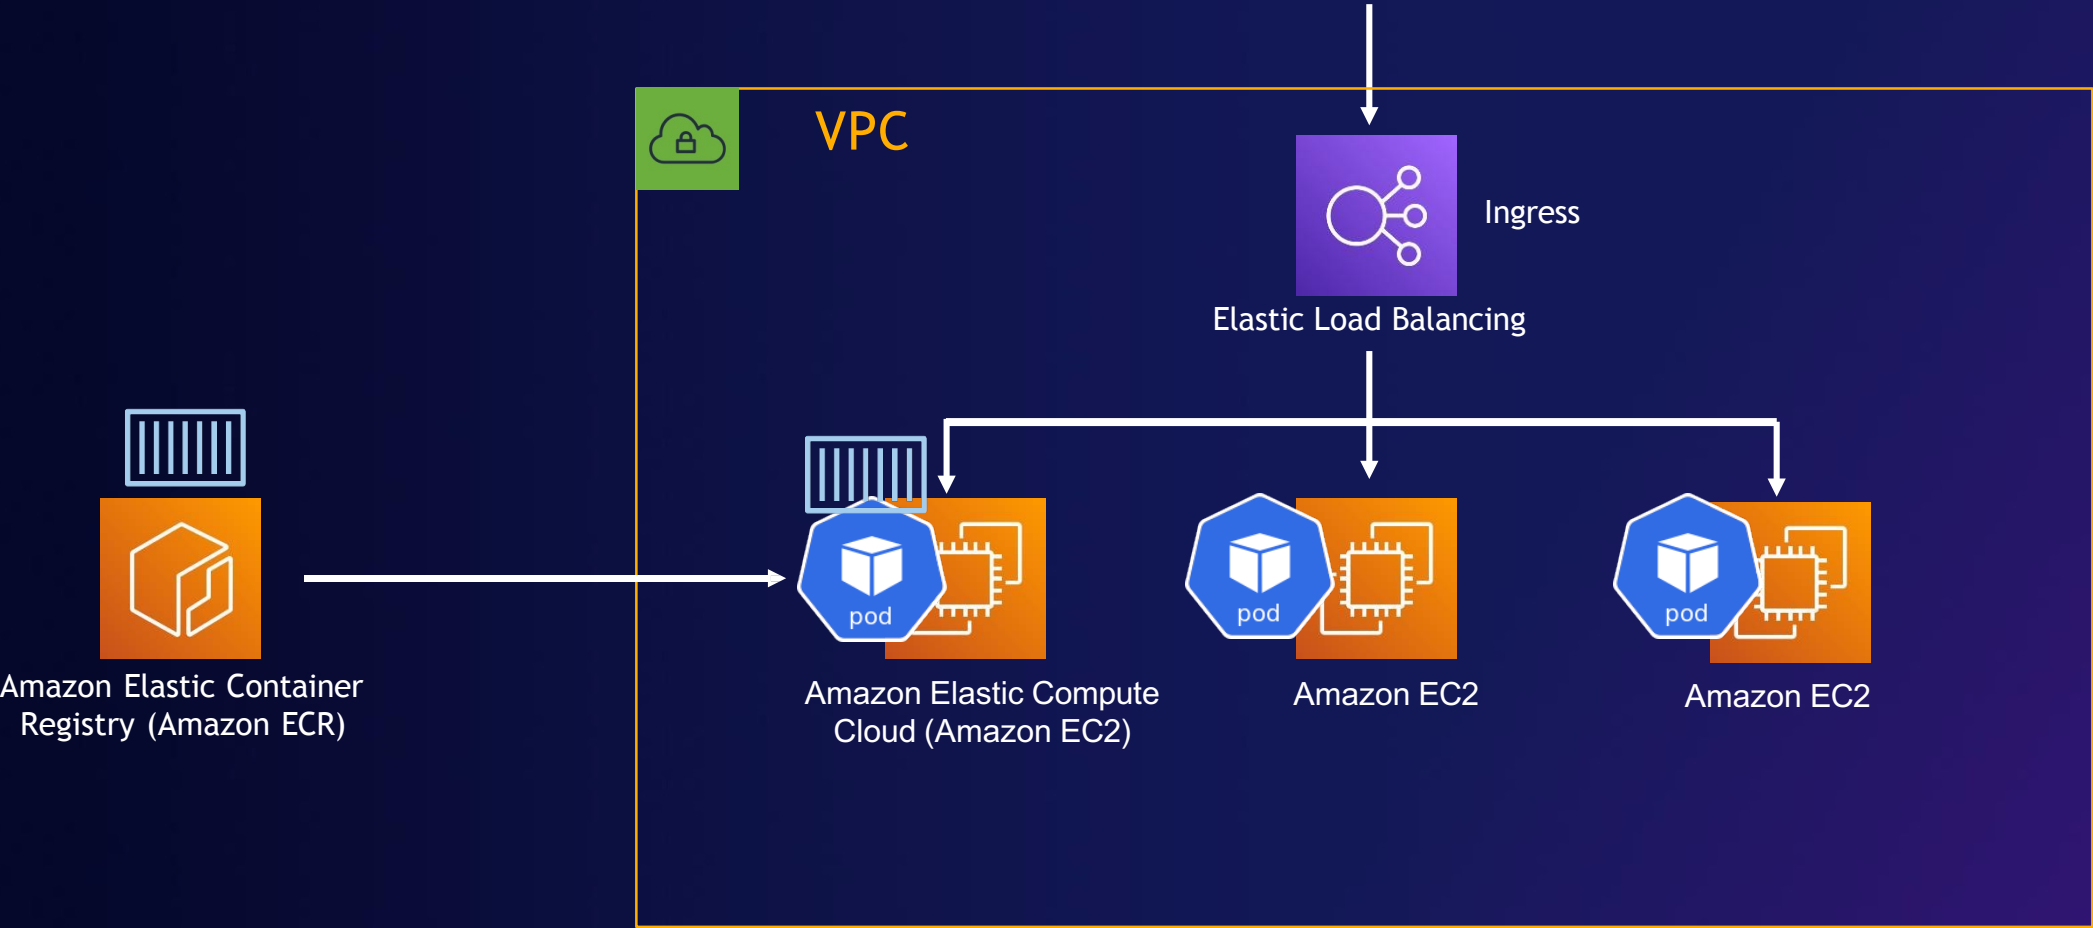
Werner Vogel
AWS CTO



Containers flow



Amazon Elastic Kubernetes
Service (Amazon EKS)



What we already know



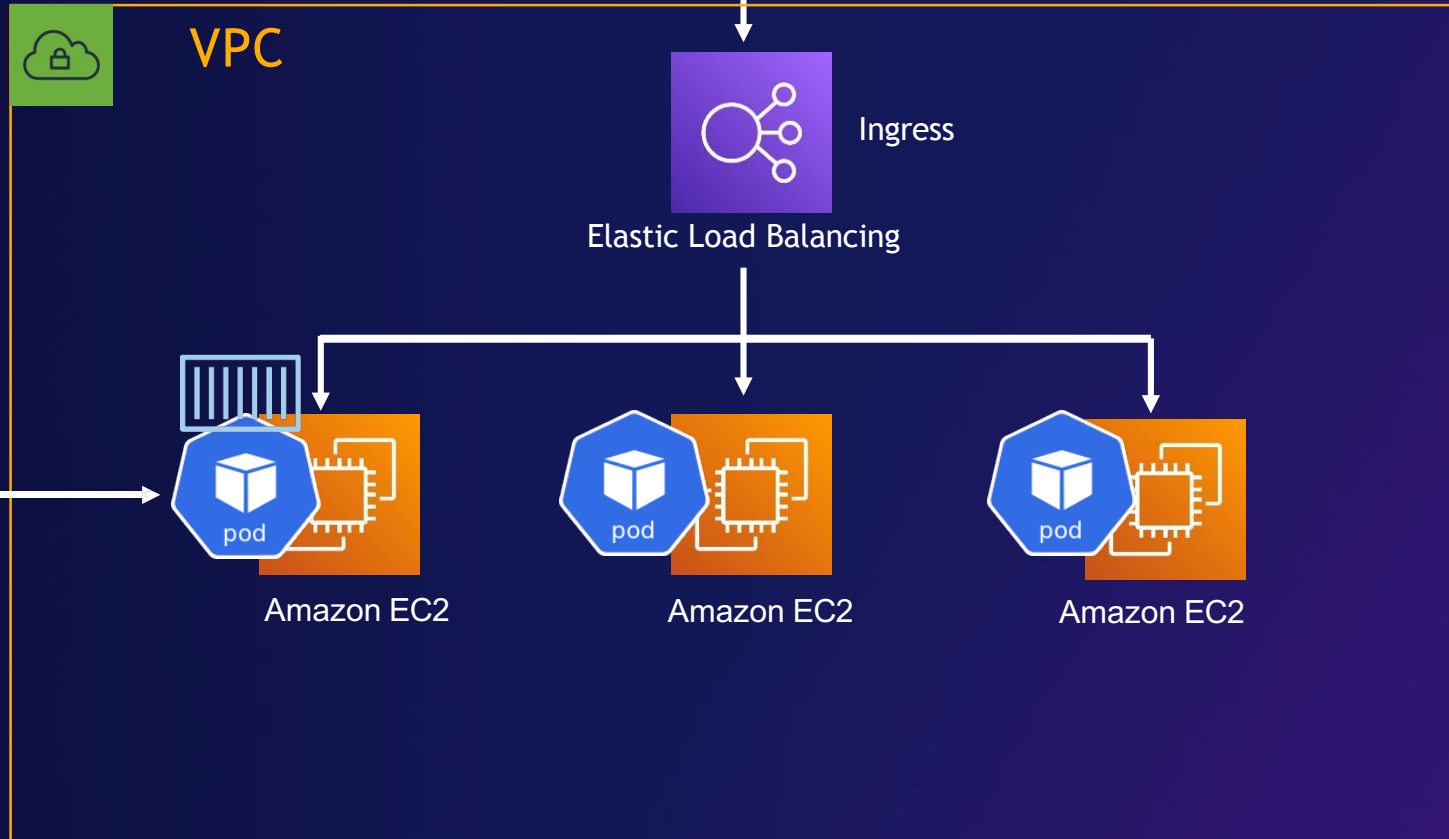
What we already know



Amazon EKS



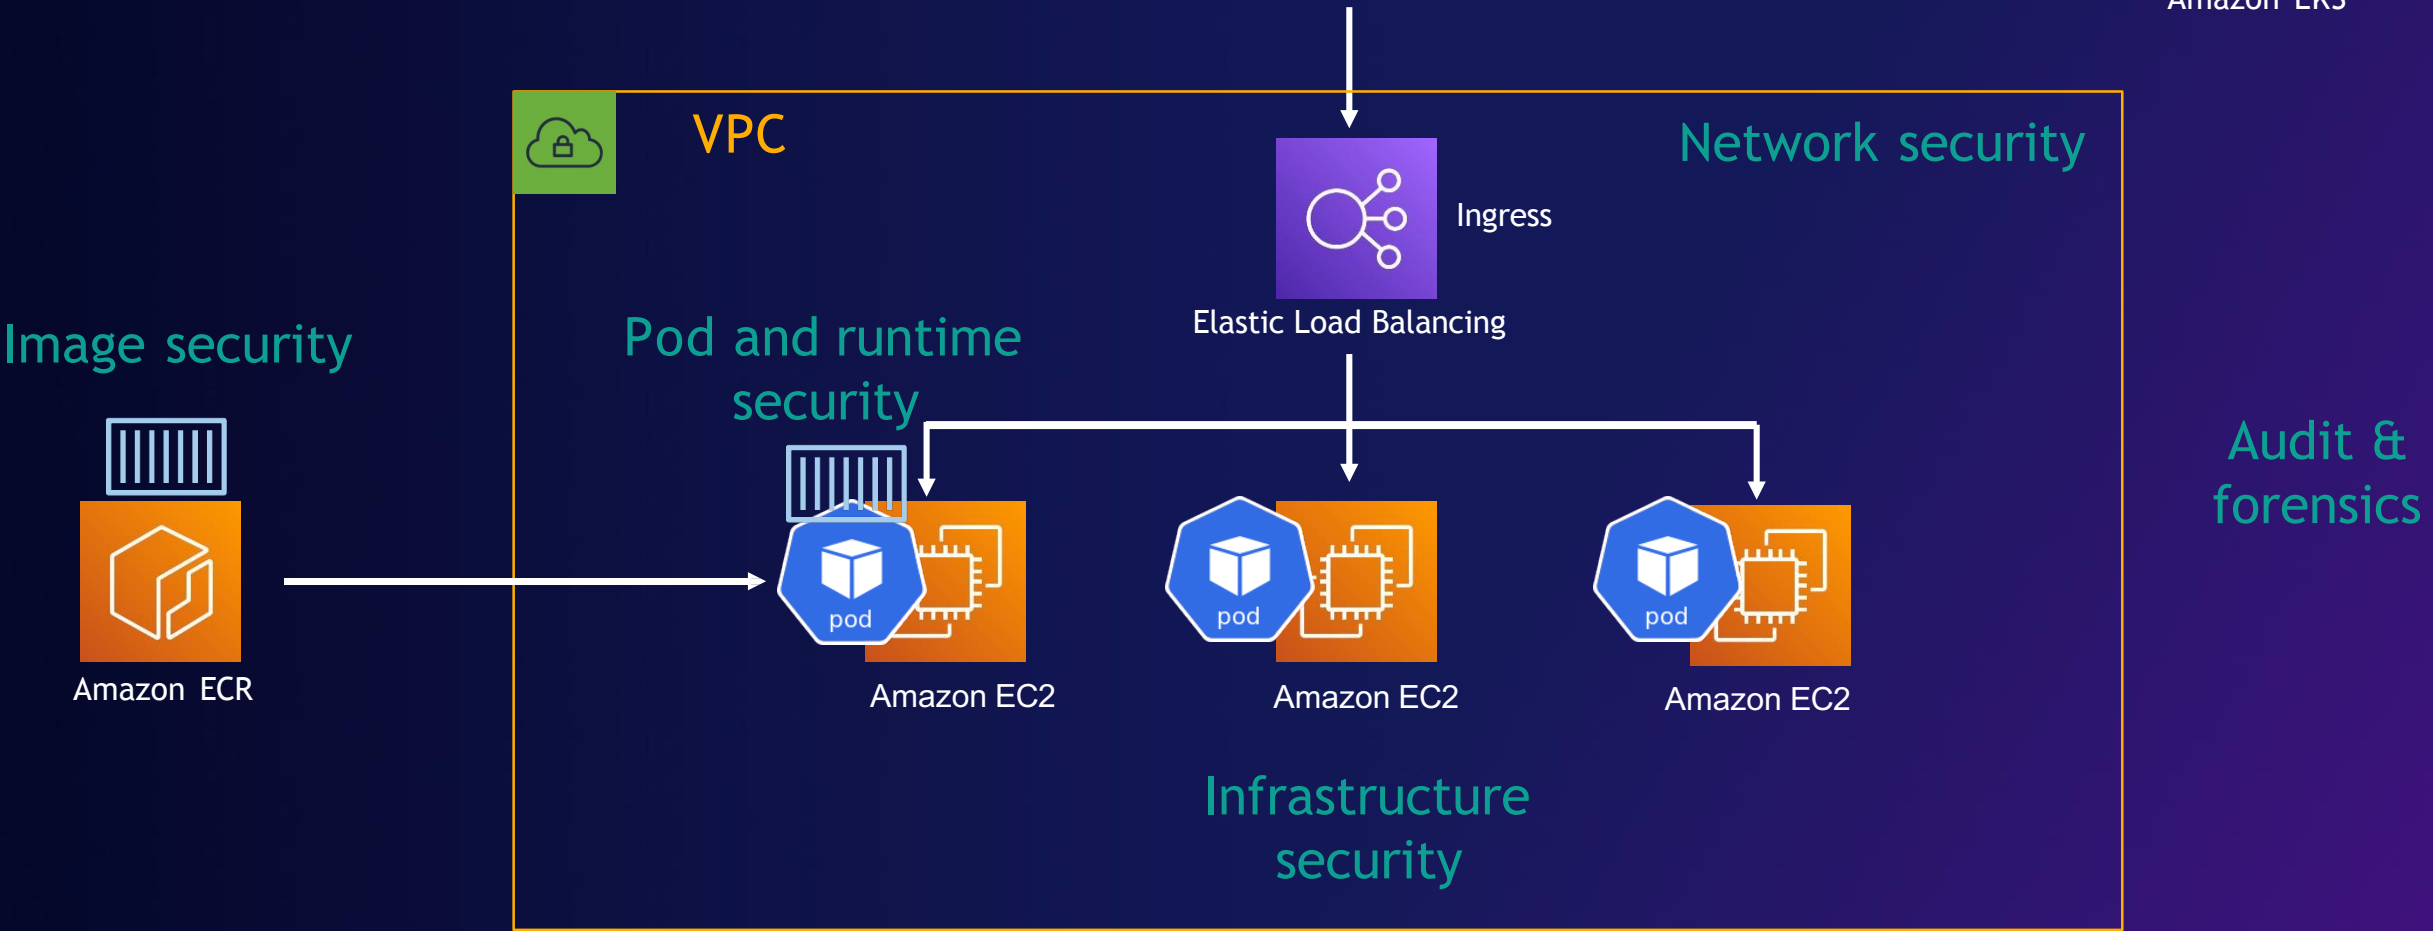
Amazon API
Gateway



Containers-specific areas



Amazon EKS





Amazon EKS

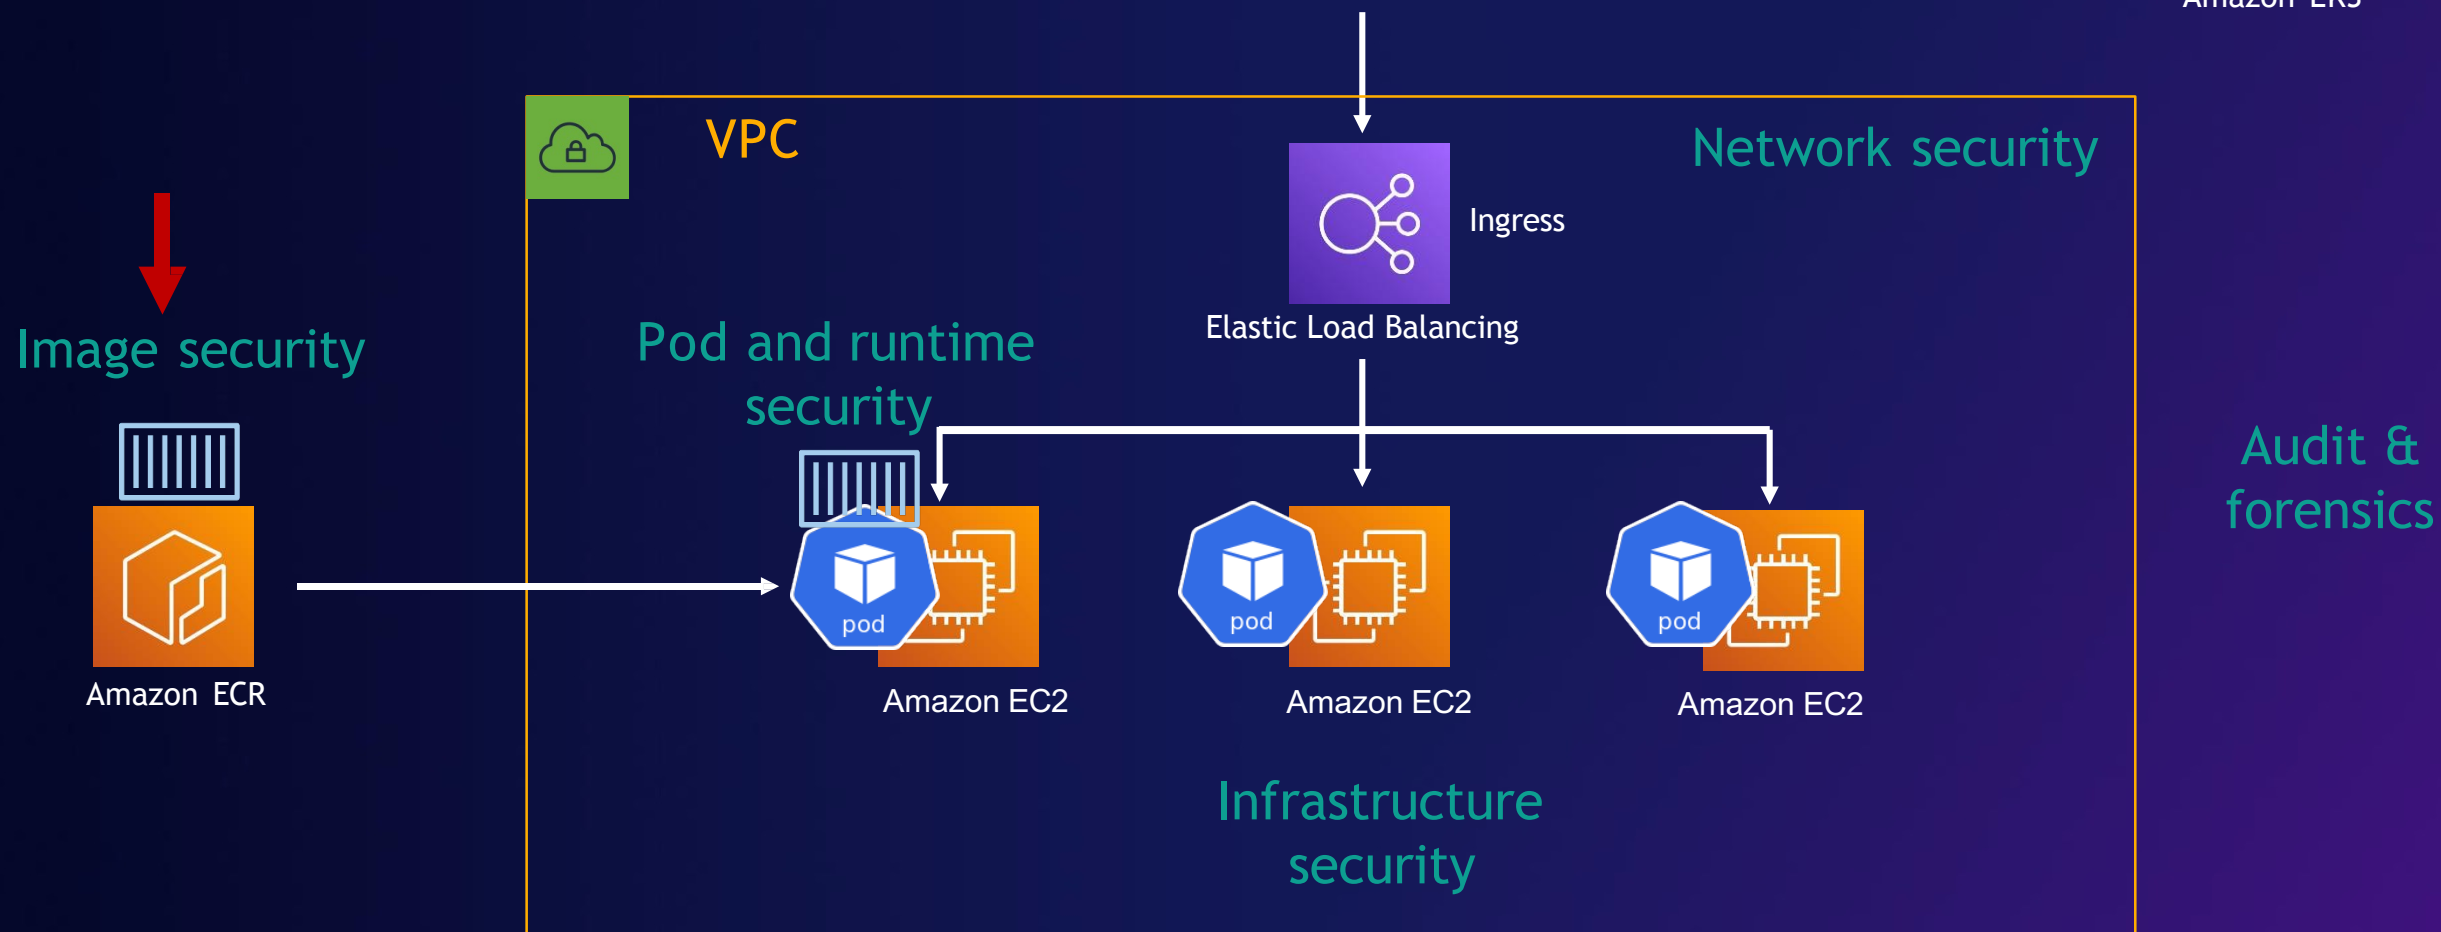



Image security

- Scan your images for vulnerabilities (CVEs)
 - ECR Basic Scanning (OS packages)
 - ECR enhanced scanning (OS and programming packages)

Inspector container image findings

Inspector > Findings > By container image > arn:aws:ecr:ap-northeast-2:628695400614:repository/spring4shell/sha256:9fd8a153a8451796c82ce0474bea5b603d5d943d172c070d1aff42e83f845d21



latest Info

Container image: sha256:9fd8a153a8451796c82ce0474bea5b603d5d943d172c070d1aff42e83f845d21

Details

AWS account

Hoseok

Repository

spring4shell

Image tags

latest

Finding summary

Package finding

22 Critical

36 High

39 Medium

By image

By layer

Findings (109)

Active

Resource ID EQUALS arn:aws:ecr:ap-northeast-2:6286954006...

Add filter

< 1 2 3 4 >

Severity	Title	Type	Age	Status
Critical	CVE-2022-1586 - pcre2	Package Vulnerability	a day	Active
Critical	CVE-2022-37434 - zlib	Package Vulnerability	19 days	Active
Critical	CVE-2022-32207 - curl, curl	Package Vulnerability	40 days	Active
Critical	CVE-2022-23219 - glibc	Package Vulnerability	83 days	Active
Critical	SNYK-JAVA-ORGCODEHAUSPLEXUS-31522 - org.codehaus.plexus:plexus-utils, org.codehaus.plexus:plexus-utils and 1 more	Package Vulnerability	7 days	Active
Critical	CVE-2021-22945 - curl, curl	Package Vulnerability	40 days	Active
Critical	GHSA-36p3-wjmg-h94x - org.springframework:spring-beans, org.springframework:spring-beans and 2 more	Package Vulnerability	83 days	Active

Image security

- Scan your images for vulnerabilities (CVEs)
 - ECR Basic Scanning (OS packages)
 - ECR enhanced scanning (OS and programming packages)
- Minimize attack surface
 - Use multi stage build
 - Create images from scratch

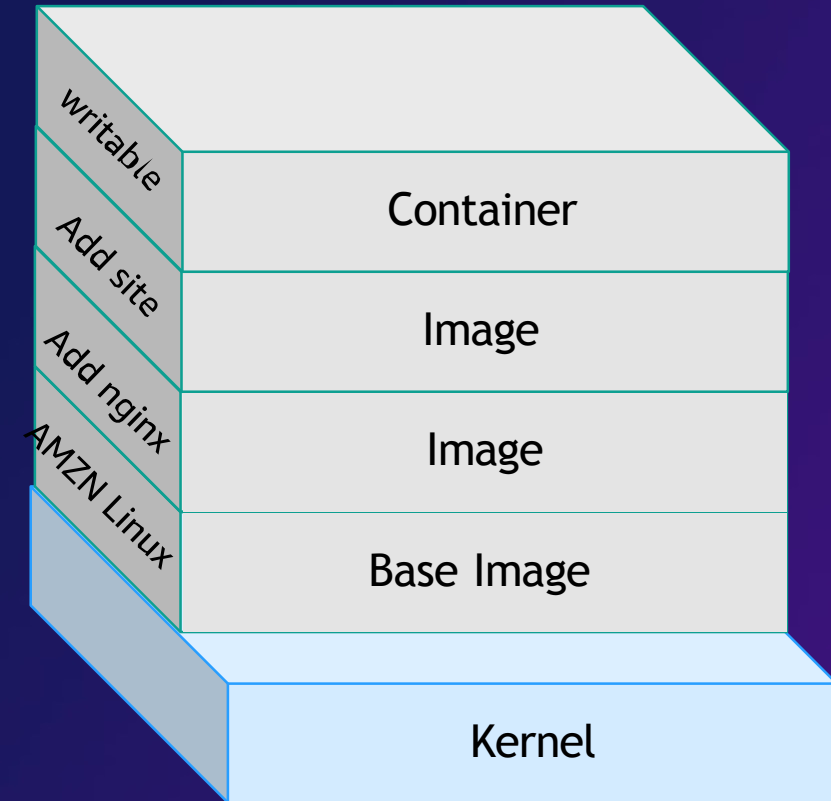
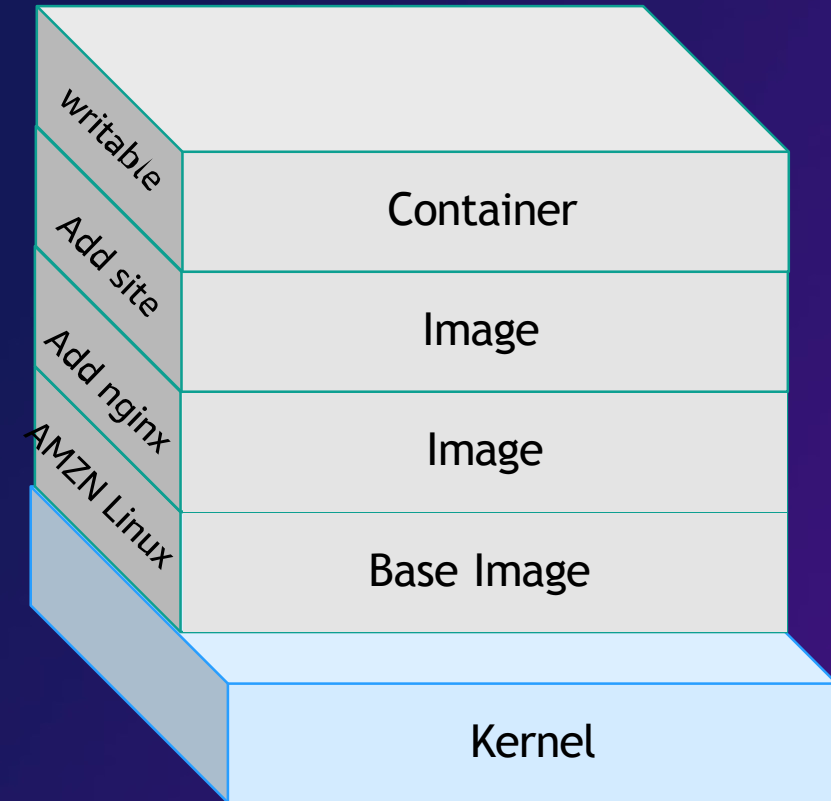


Image security

- Scan your images for vulnerabilities (CVEs)
 - ECR Basic Scanning (OS packages)
 - ECR enhanced scanning (OS and programming packages)
- Minimize attack surface
 - Use multi stage build
 - Create images from scratch
- Run the application as a non-root user
 - Lint your Dockerfiles with [Dockle](#) or [Hadolint](#)



Lint Dockerfiles

```
FROM alpine:3.16
```

```
LABEL maintainer="NGINX Docker Maintainers <docker-maint@nginx.com>"
```

```
ENV NGINX_VERSION 1.22.0
```

```
ENV PKG_RELEASE 1
```

```
RUN set -x \
```

```
# create nginx user/group first, to be consistent throughout docker variants
```

```
&& addgroup -g 101 -S nginx \
```

```
&& adduser -S -D -H -u 101 -h /var/cache/nginx -s /sbin/nologin -G nginx -g nginx nginx \
```

```
&& apkArch="$(cat /etc/apk/arch)" \
```

```
&& nginxPackages=" \
```

```
    nginx=${NGINX_VERSION}-r${PKG_RELEASE} \
```

```
" \
```

```
# install prerequisites for public key and pkg-oss checks
```

```
&& apk add --no-cache --virtual .checksum-deps \
```

```
    openssl \
```

```
&& case "$apkArch" in \
```

```
    x86_64|aarch64) \
```

```
# arches officially built by upstream
```

```
set -x \
```

```
&& KEY_SHA512="e7fa8303923d9b95db37a77ad46c68fd4755ff935d0a534d26eba83de193c76166c68bfe7f65471bf8881004ef4aa6df3e34689c305662750c0172fca5d8552a *stdin" \
```

```
&& wget -O /tmp/nginx_signing.rsa.pub https://nginx.org/keys/nginx_signing.rsa.pub \
```

```
&& if [ "$(openssl rsa -pubin -in /tmp/nginx_signing.rsa.pub -text -noout | openssl sha512 -r)" = "$KEY_SHA512" ]; then \
```

```
    echo "key verification succeeded!"; \
```

```
    mv /tmp/nginx_signing.rsa.pub /etc/apk/keys/; \
```

```
else \
```

```
    echo "key verification failed!"; \
```

```
    exit 1; \
```

```
fi \
```

```
&& apk add -X "https://nginx.org/packages/alpine/v$(egrep -o '^[0-9]+\.[0-9]+' /etc/alpine-release)/main" --no-cache $nginxPackages \
```

```
;; \
```

```
*) \
```

```
[ec2-user@ip-10-0-1-211 ~]$ ./dockle alpine:3.16
```

```
WARN - CIS-DI-0001: Create a user for the container
```

```
* Last user should not be root
```

```
INFO - CIS-DI-0005: Enable Content trust for Docker
```

```
* export DOCKER_CONTENT_TRUST=1 before docker pull/build
```

```
INFO - CIS-DI-0006: Add HEALTHCHECK instruction to the container image
```

```
* not found HEALTHCHECK statement
```

```
[ec2-user@ip-10-0-1-211 ~]$
```

```
[ec2-user@ip-10-0-1-211 ~]$ hadolint Dockerfile
```

```
Dockerfile:13 DL3018 warning: Pin versions in apk add. Instead of `apk add <package>` use `apk add <package>=<version>`
```

```
Dockerfile:13 DL3047 info: Avoid use of wget without progress bar. Use `wget --progress=dot:giga <url>`.Or consider using `-q` or `-nv` (shorthands for `--quiet` or `--no-verbose`).
```

```
Dockerfile:13 SC2086 info: Double quote to prevent globbing and word splitting.
```

```
Dockerfile:13 DL4006 warning: Set the SHELL option -o pipefail before RUN with a pipe in it. If you are using /bin/sh in an alpine image or if your shell is symlinked to busybox then consider explicitly setting your SHELL to /bin/ash, or disable this check
```

Image security

- Scan your images for vulnerabilities (CVEs)
 - ECR Basic Scanning (OS packages)
 - ECR enhanced scanning (OS and programming packages)
- Minimize attack surface
 - Use multi stage build
 - Create images from scratch
- Run the application as a non-root user
 - Lint your Dockerfiles with [Dockle](#) or [Hadolint](#)
- “Defang” your containers
 - Remove files with the SETUID and SETGID bits from image

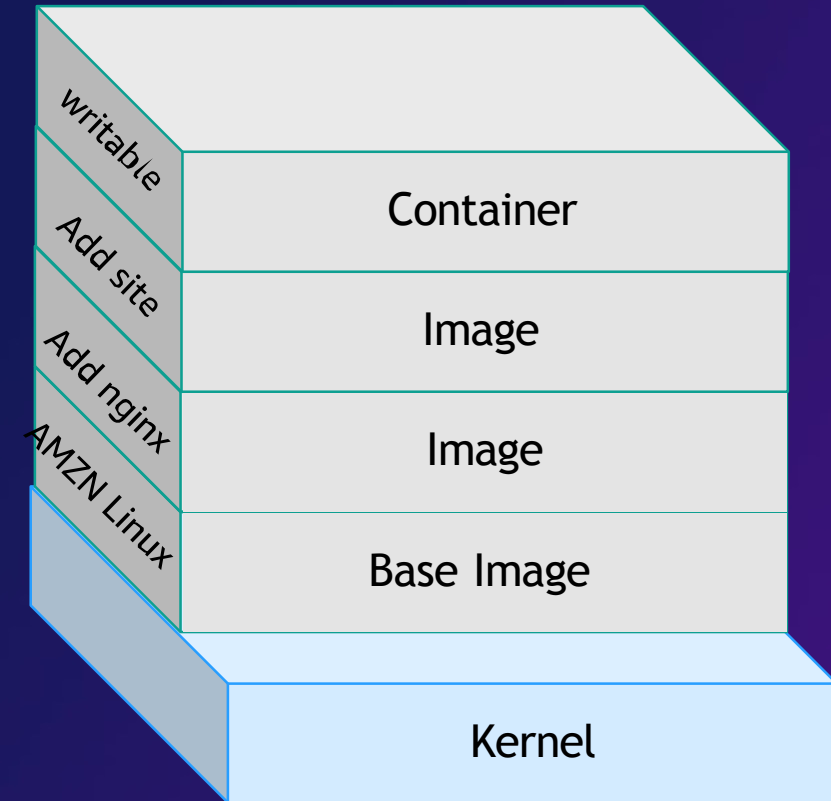


Image security - dive

De-fang your images

- Remove files with the SETUID and SETGID bits from the image
- Dive (<https://github.com/wagoodman/dive>)

```
1/1 + [Icons] Tilix: Default
1: Terminal ▾
[● Layers]
Cmp Image ID Size Command
sha256:cd7100a72410606589 4.1 MB FROM sha256:cd7100a72410606589
sha256:f03b1ccbacace8c82e 2.1 kB #(nop) ADD file:63d4894bd0857354b
sha256:d2a26525cdac6c4358 0 B mkdir /root/example
sha256:9f50561518484bb62a 2.1 kB cp /somefile.txt /root/example/somefile1.txt
sha256:edf209d09263ab03c7 2.1 kB cp /somefile.txt /root/example/somefile2.txt
sha256:bb70c7aab83602e8a6 2.1 kB cp /somefile.txt /root/example/somefile3.txt
sha256:b0a712d3d00bc83c48 2.1 kB mv /root/example/somefile3.txt /root/example/
sha256:05e8660b8f90e77b93 0 B rm -rf /root/example/

[Current Layer Contents]
Permission UID:GID Size Filetree
drwxr-xr-x 0:0 805 kB bin
drwxr-xr-x 0:0 0 B dev
drwxr-xr-x 0:0 251 kB etc
drwxr-xr-x 0:0 0 B home
drwxr-xr-x 0:0 2.7 MB lib
drwxr-xr-x 0:0 0 B media
drwxr-xr-x 0:0 0 B cdrom
drwxr-xr-x 0:0 0 B floppy
drwxr-xr-x 0:0 0 B usb
drwxr-xr-x 0:0 0 B mnt
dr-xr-xr-x 0:0 0 B proc
drwx----- 0:0 6.2 kB root
drwxr-xr-x 0:0 6.2 kB example
-rw-r--r-- 0:0 2.1 kB somefile1.txt
-rw-r--r-- 0:0 2.1 kB somefile2.txt
-rw-r--r-- 0:0 2.1 kB somefile3.txt
drwxr-xr-x 0:0 0 B run
drwxr-xr-x 0:0 213 kB sbin
-rw-rw-r-- 0:0 2.1 kB somefile.txt
drwxr-xr-x 0:0 0 B srv
drwxr-xr-x 0:0 0 B sys
drwxrwxrwx 0:0 0 B tmp
drwxr-xr-x 0:0 176 kB usr
drwxr-xr-x 0:0 0 B var

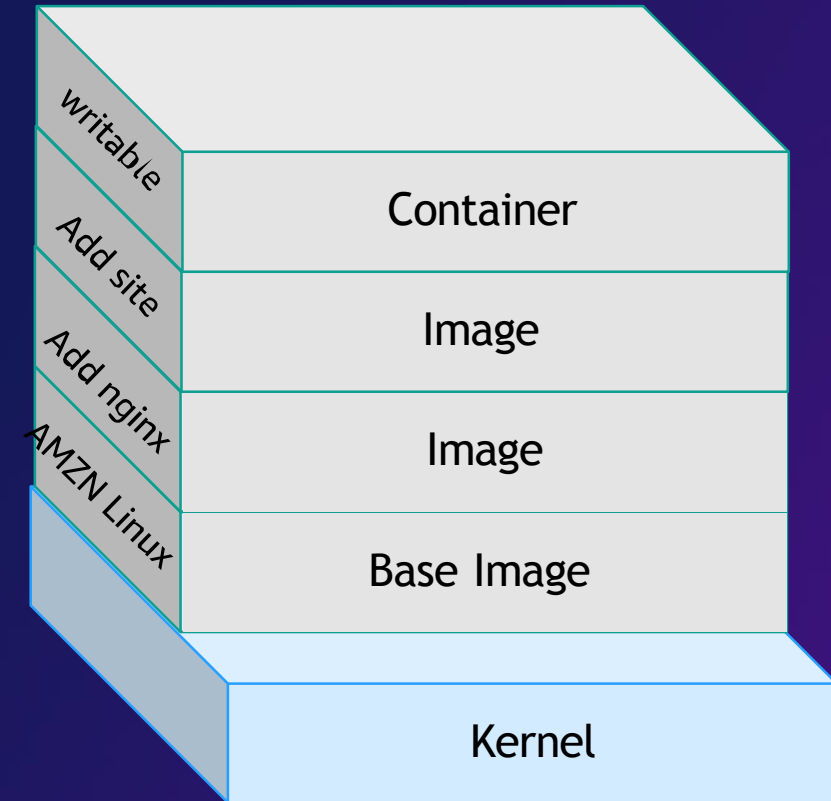
[Image & Layer Details]
Layer Command
/bin/sh -c cp /somefile.txt /root/example/somefile3.txt

Image efficiency score: 99 %
Potential wasted space: 6.2 kB

Count Total Space Path
2 4.2 kB /root/example
2 2.1 kB /root/example/somefile3.txt
```

Image security

- Scan your images for vulnerabilities (CVEs)
 - ECR Basic Scanning (OS packages)
 - ECR enhanced scanning (OS and programming packages)
- Minimize attack surface
 - Use multi stage build
 - Create images from scratch
- Run the application as a non-root user
 - Lint your Dockerfiles with [Dockle](#) or [Hadolint](#)
- “Defang” your containers
 - Remove files with the SETUID and SETGID bits from image
- Use endpoint policies and private endpoints with Elastic Container Registry



VPC endpoint policy

VPC > Endpoints > Create endpoint

Create endpoint [Info](#)

There are three types of VPC endpoints – Interface endpoints, Gateway Load Balancer endpoints, and Gateway endpoints. Interface endpoints and Gateway Load Balancer endpoints are powered by AWS PrivateLink, and use an Elastic Network Interface (ENI) as an entry point for traffic destined to the service. Interface endpoints are typically accessed using the public or private DNS name associated with the service, while Gateway endpoints and Gateway Load Balancer endpoints serve as a target for a route in your route table for traffic destined for the service.

Endpoint settings

Name tag - optional

Creates a tag with a key of 'Name' and a value that you specify.

Service category

Select the service category

☒ AWS services

Services provided by Amazon

☐ PrivateLink Ready partner services

Services with an AWS Service Ready designation

☐ AWS Marketplace services

Services that you've purchased through AWS Marketplace

☐ Other endpoint services

Find services shared with you by service name

Services (1/5)

search: ecr. X

search: com.amazonaws.ap-northeast-2.s3 X

Clear filters

	Service Name	Owner	Type
<input type="radio"/>	com.amazonaws.ap-northeast-2.ecr.api	amazon	Interface
<input type="radio"/>	com.amazonaws.ap-northeast-2.ecr.dkr	amazon	Interface
<input type="radio"/>	com.amazonaws.ap-northeast-2.s3	amazon	Gateway
<input type="radio"/>	com.amazonaws.ap-northeast-2.s3	amazon	Interface
<input type="radio"/>	com.amazonaws.ap-northeast-2.s3-outposts	amazon	Interface

Create an endpoint policy for your Amazon ECR VPC endpoints

A VPC endpoint policy is an IAM resource policy that you attach to an endpoint when you create or modify the endpoint. If you don't attach a policy when you create an endpoint, AWS attaches a default policy for you that allows full access to the service. An endpoint policy doesn't override or replace IAM user policies or service-specific policies. It's a separate policy for controlling access from the endpoint to the specified service. Endpoint policies must be written in JSON format. For more information, see [Controlling Access to Services with VPC Endpoints](#) in the *Amazon VPC User Guide*.

We recommend creating a single IAM resource policy and attaching it to both of the Amazon ECR VPC endpoints.

The following is an example of an endpoint policy for Amazon ECR. This policy enables a specific IAM role to pull images from Amazon ECR.

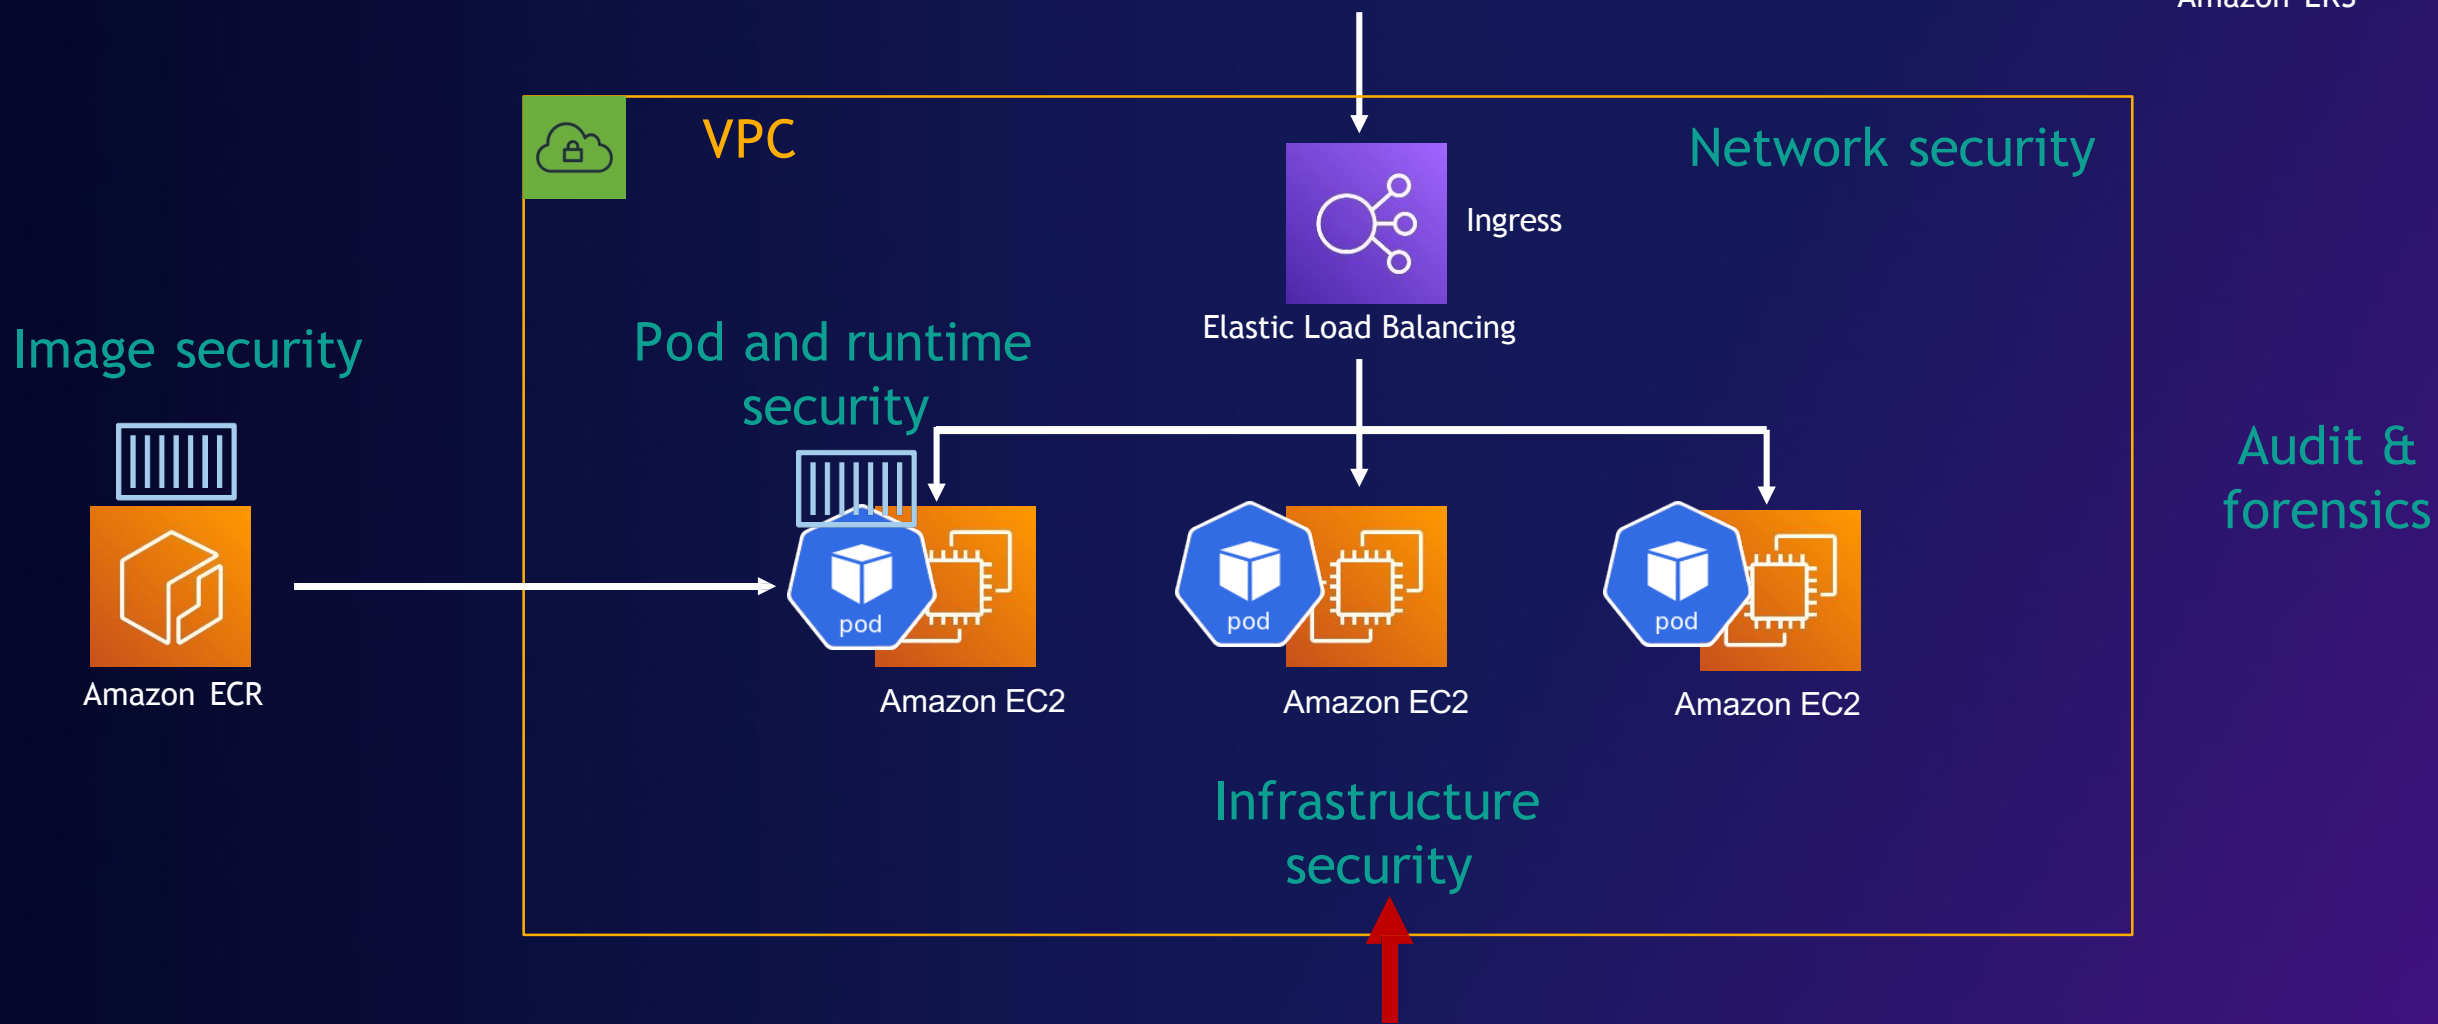
```
{
  "Statement": [{
    "Sid": "AllowPull",
    "Principal": {
      "AWS": "arn:aws:iam::1234567890:role/role_name"
    },
    "Action": [
      "ecr:BatchGetImage",
      "ecr:GetDownloadUrlForLayer",
      "ecr:GetAuthorizationToken"
    ],
    "Effect": "Allow",
    "Resource": "*"
  }]
}
```

The following endpoint policy example prevents a specified repository from being deleted.

```
{
  "Statement": [{
    "Sid": "AllowAll",
    "Principal": "*",
    "Action": "*",
    "Effect": "Allow",
    "Resource": "*"
  }, {
    "Sid": "PreventDelete",
    "Principal": "*",
    "Action": "ecr:DeleteRepository",
    "Effect": "Deny",
    "Resource": "arn:aws:ecr:region:1234567890:repository/repository_name"
  }]
}
```




Amazon EKS




Infrastructure security

- Use an OS that is optimized for running containers
 - Amazon EKS-optimized AL2 or [Bottlerocket](#)
- Upgrade to latest AMIs
 - Use Managed Node Group for rolling deployment


Infrastructure security - AMI

amazon-eks-node-1.22-v20220 (4 filtered, 16 unfiltered)



Community AMIs

Community AMIs contain all AMIs that are public, therefore anyone can publish an AMI and it will show in this catalog. This catalog can also contain paid products. When using community AMIs it is best practice to ensure you know and trust the publisher before launching an AMI.



Verified provider

amazon-eks-node-1.22-v20220824

ami-03f5c5b7ba2015719

EKS Kubernetes Worker AMI with AmazonLinux2 image, (k8s: 1.22.12, docker: 20.10.17-1.amzn2, containerd: 1.6.6-1.amzn2)

Platform: Other Linux

Architecture: x86_64

Owner: 602401143452


Publish date: 2022-08-24

Root device type: ebs

Virtualization: hvm

ENA enabled: Yes

Select



Verified provider

amazon-eks-node-1.22-v20220802

ami-09e0355345619d14a

EKS Kubernetes Worker AMI with AmazonLinux2 image, (k8s: 1.22.9, docker: 20.10.13-2.amzn2, containerd: 1.4.13-3.amzn2)

Platform: Other Linux

Architecture: x86_64

Owner: 602401143452


Publish date: 2022-08-02

Root device type: ebs

Virtualization: hvm

ENA enabled: Yes

Select



Verified provider

amazon-eks-node-1.22-v20220811

ami-028f26a5a0798bf1e

EKS Kubernetes Worker AMI with AmazonLinux2 image, (k8s: 1.22.12, docker: 20.10.13-2.amzn2, containerd: 1.4.13-3.amzn2)

Platform: Other Linux

Architecture: x86_64

Owner: 602401143452


Publish date: 2022-08-11

Root device type: ebs

Virtualization: hvm

ENA enabled: Yes

Select



Verified provider

amazon-eks-node-1.22-v20220914

ami-0167f7c2f2c676793

EKS Kubernetes Worker AMI with AmazonLinux2 image, (k8s: 1.22.12, docker: 20.10.17-1.amzn2, containerd: 1.6.6-1.amzn2)

Platform: Other Linux

Architecture: x86_64

Owner: 602401143452

Publish date: 2022-09-14

Root device type: ebs


Virtualization: hvm

ENA enabled: Yes

Select

AMI Release v20220914

Latest

 ljosyula released this 3 days ago · 4 commits to master since this release · v20220914 · 802add4

AMI Release v20220914

- amazon-eks-gpu-node-1.23-v20220914
- amazon-eks-gpu-node-1.22-v20220914
- amazon-eks-gpu-node-1.21-v20220914
- amazon-eks-gpu-node-1.20-v20220914
- amazon-eks-arm64-node-1.23-v20220914
- amazon-eks-arm64-node-1.22-v20220914
- amazon-eks-arm64-node-1.21-v20220914
- amazon-eks-arm64-node-1.20-v20220914
- amazon-eks-node-1.23-v20220914
- amazon-eks-node-1.22-v20220914
- amazon-eks-node-1.21-v20220914
- amazon-eks-node-1.20-v20220914

Release versions for these AMIs:

- 1.23.9-20220914
- 1.22.12-20220914
- 1.21.14-20220914
- 1.20.15-20220914

Binaries used to build these AMIs are published:

- s3://amazon-eks/1.23.9/2022-07-27/
- s3://amazon-eks/1.22.12/2022-07-27/
- s3://amazon-eks/1.21.14/2022-07-27/
- s3://amazon-eks/1.20.15/2022-07-27/

AMI details:

- kernel: 5.4.209-116.367.amzn2
- dockerd: 20.10.17-1.amzn2
- containerd: 1.6.6-1.amzn2
- runc: 1.1.3-1.amzn2
- cuda: 470.57.02-1
- nvidia-container-runtime-hook: 1.4.0-1.amzn2
- SSM agent: 3.1.1732.0-1.amzn2



Infrastructure security - Managed Node Group

eks-121-proxy02

Delete cluster

Your current user or role does not have access to Kubernetes objects on this EKS cluster
This may be due to the current user or role not having Kubernetes RBAC permissions to describe cluster resources or not having an entry in the cluster's auth config map. [Learn more](#)

New Kubernetes versions are available for this cluster. [Learn more](#)

Update now

▼ Cluster info [Info](#)

Kubernetes version [Info](#)
1.21

Status
✔ Active

Provider
EKS

Overview

Resources

Compute

Networking

Add-ons

Authentication

Logging

Update history

Tags

Nodes (0) [Info](#)

🔍 Filter Nodes by property or value

< 1 >

Node name	Instance type	Node group	Created	Status
No Nodes This cluster does not have any Nodes, or you don't have permission to view them.				

Node groups (1) [Info](#)

Edit

Delete

Add node group

Group name	Desired size	AMI release version	Launch template	Status
<div></div> managed-ng01	0	1.21.12-20220629 Update now	eksctl-eks-121-proxy02-nodegroup-managed-ng01 (1)	✔ Active

Update node group version: managed-ng01

☒ Update node group version
Update your node group AMI release version to the latest available version for your 1.21 cluster.

☐ Change launch template version
Change your node group to use a different launch template version.

Update strategy

Select how to update this node group.

Rolling update

This option respects pod disruption budgets for your cluster. The update fails if Amazon EKS is unable to gracefully drain the pods that are running on this node group due to a pod disruption budget issue.

[Learn more](#)

[AMI release version notes](#)

Cancel

Update

Infrastructure security

- Use an OS that is optimized for running containers
 - Amazon EKS-optimized AL2 or [Bottlerocket](#)
- Upgrade to latest AMIs
 - Use Managed Node Group for rolling deployment
- Deploy workers onto private subnets

Infrastructure security

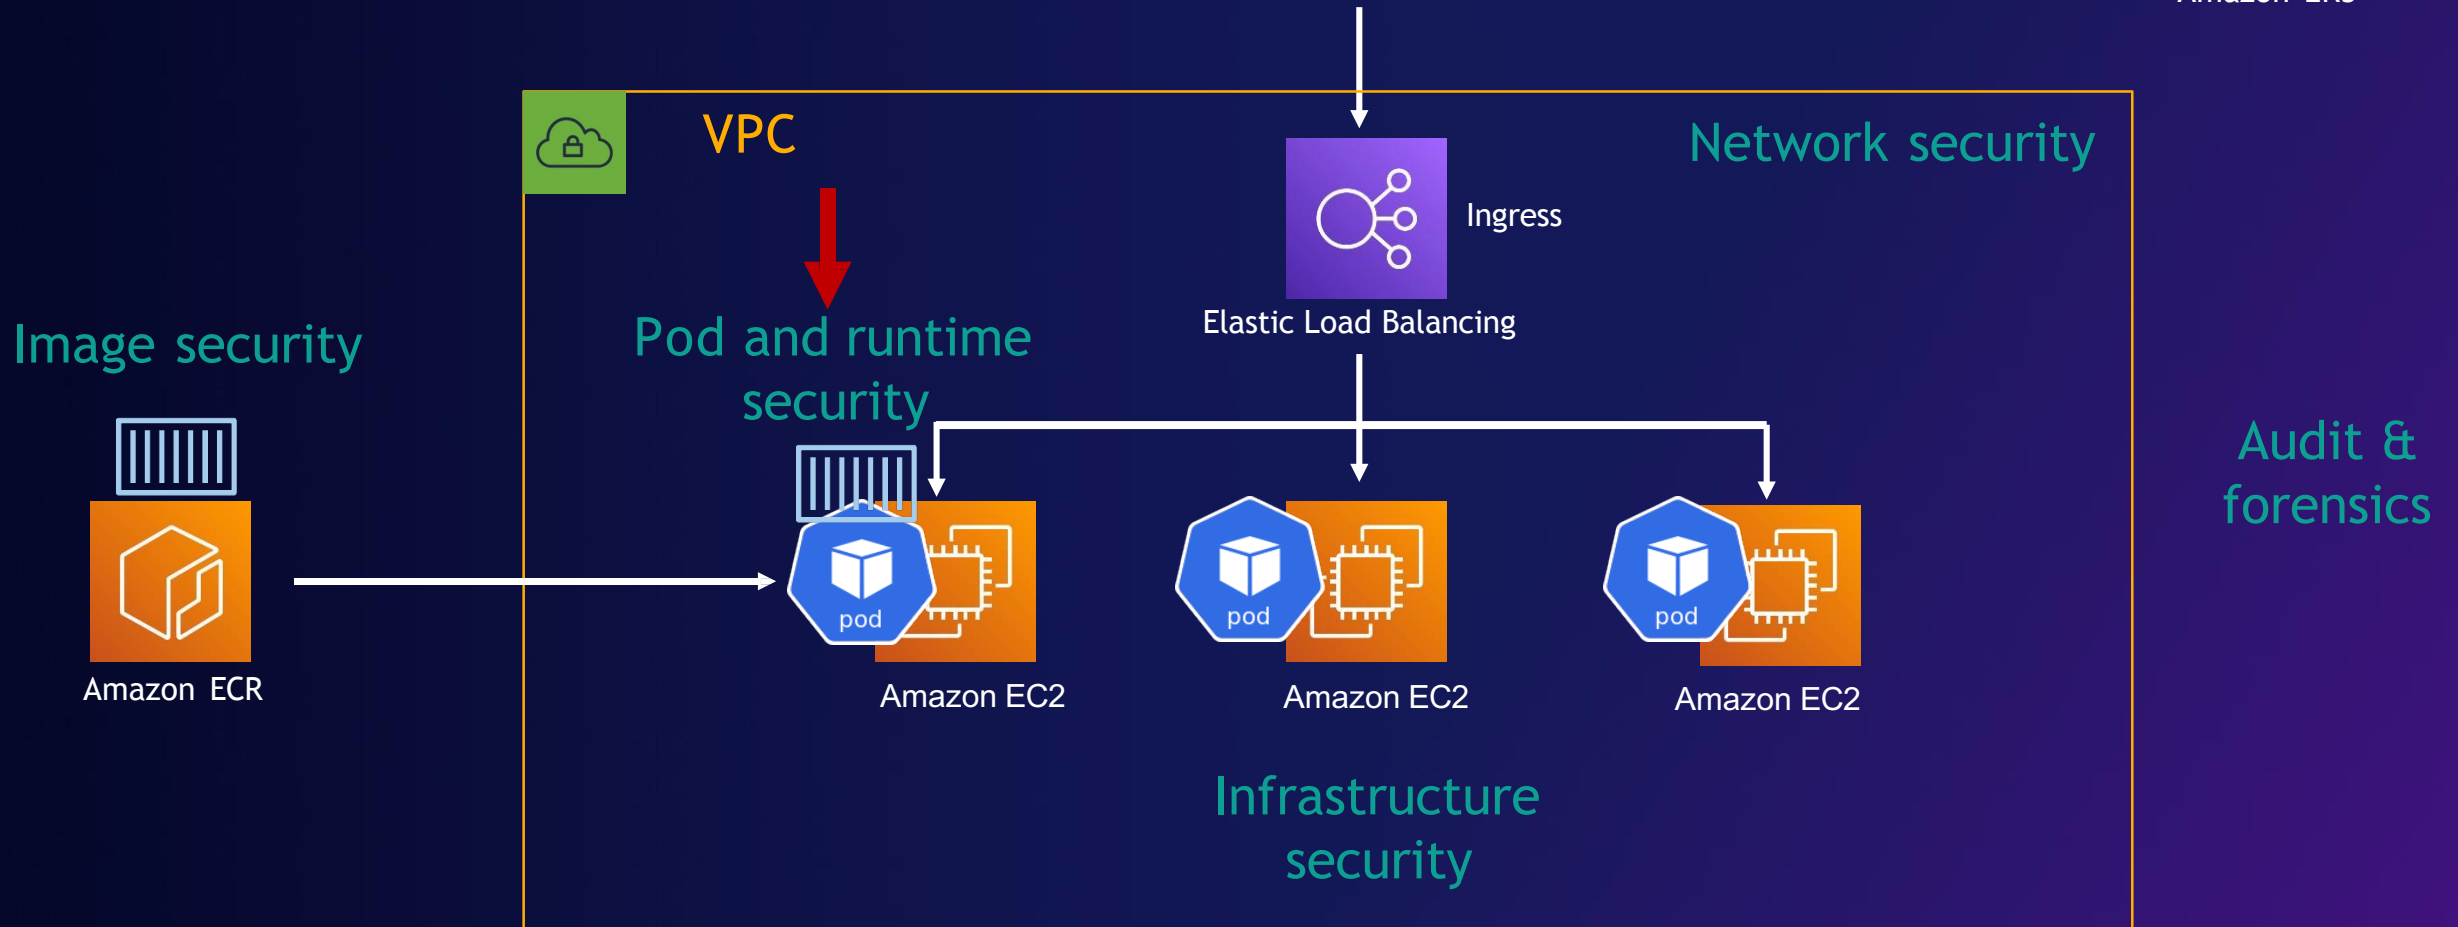
- Use an OS that is optimized for running containers
 - Amazon EKS-optimized AL2 or [Bottlerocket](#)
- Upgrade to latest AMIs
 - Use Managed Node Group for rolling deployment
- Deploy workers onto private subnets
- Minimize and audit host access
 - AWS Systems Manager, Session Manager

Infrastructure security

- Use an OS that is optimized for running containers
 - Amazon EKS-optimized AL2 or [Bottlerocket](#)
- Upgrade to latest AMIs
 - Use Managed Node Group for rolling deployment
- Deploy workers onto private subnets
- Minimize and audit host access
 - AWS Systems Manager, Session Manager
- Run kube-bench to continually assess alignment with Amazon EKS CIS benchmarks

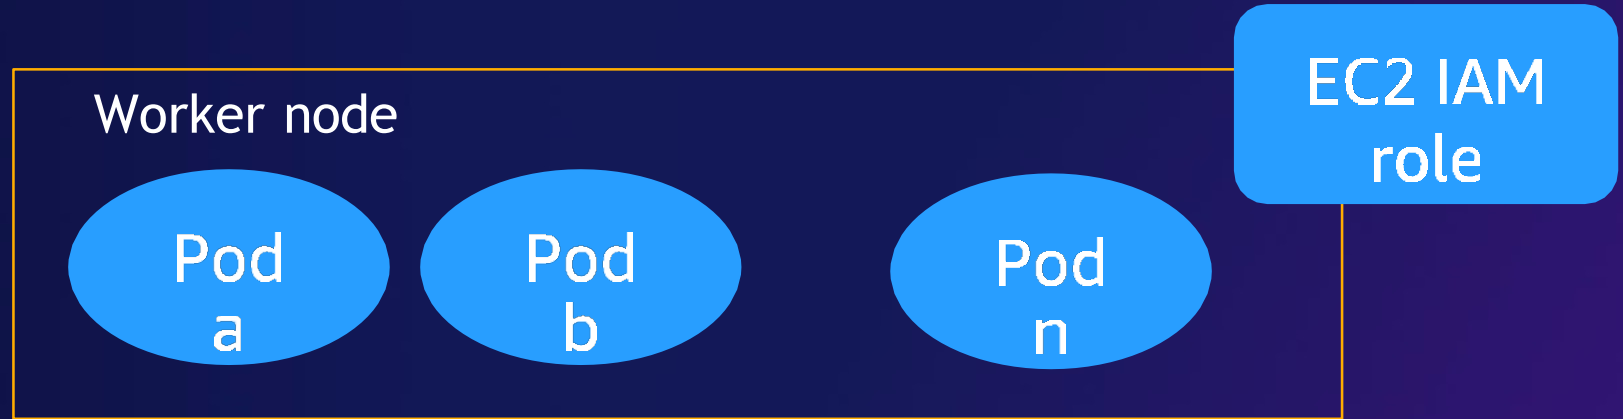


Amazon EKS

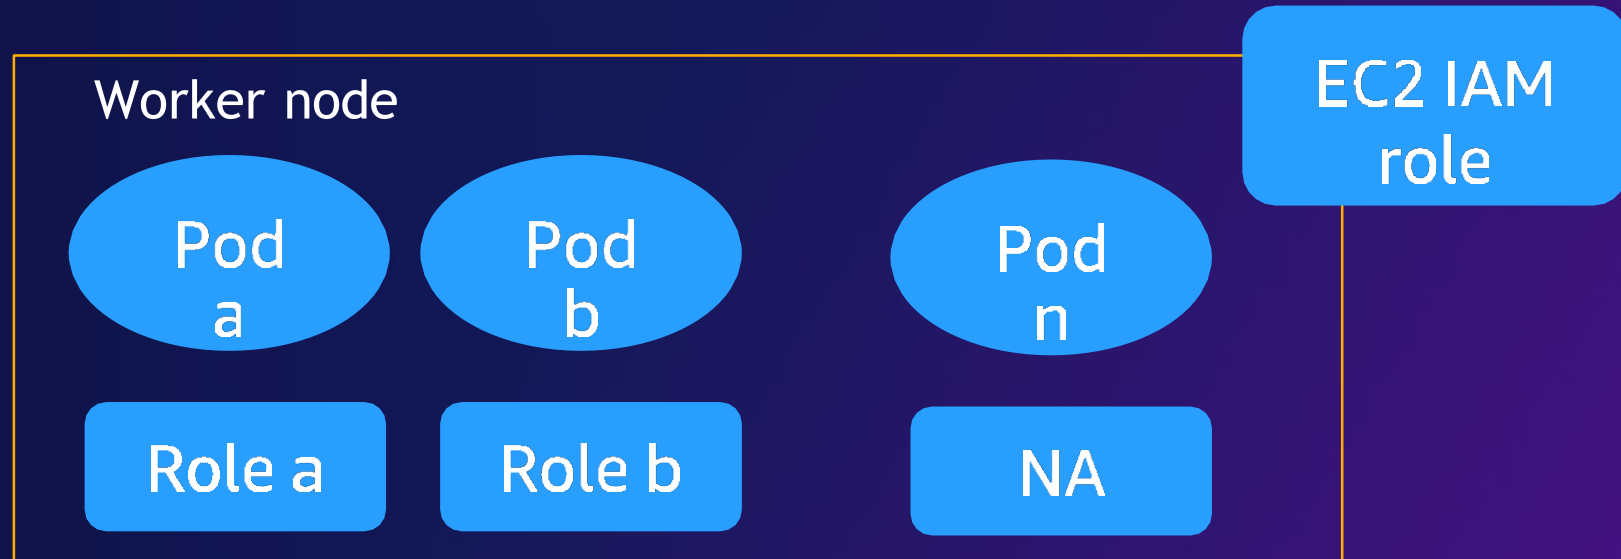


IAM roles for service accounts

No IRSA



IRSA configured

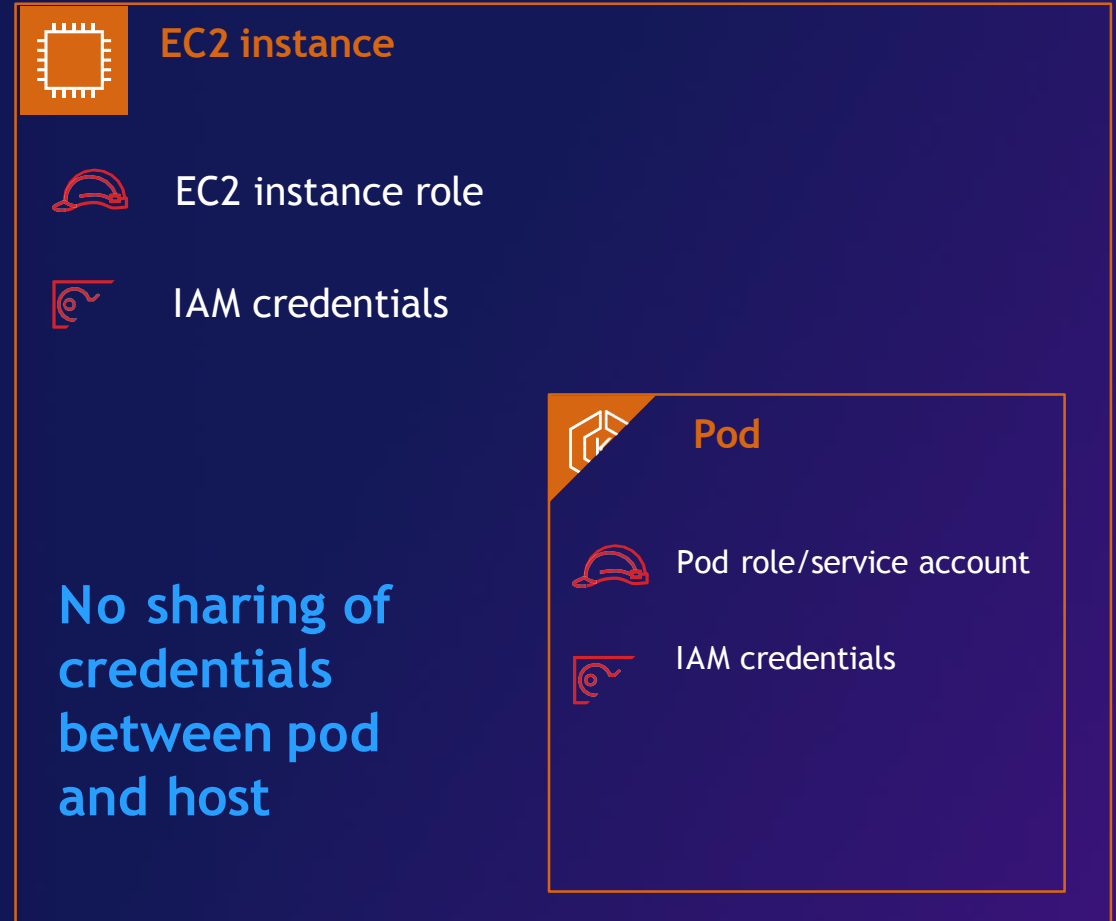


Use IAM roles for service accounts

Map IAM roles to Kubernetes service accounts

Enable least privilege for Kubernetes pods

Natively supported by Amazon EKS and available as open source



Pod and runtime security

- Implement IRSA (IAM roles for service accounts)
- Run dynamic scans - Falco, Prisma Twistlock, Aqua, etc.

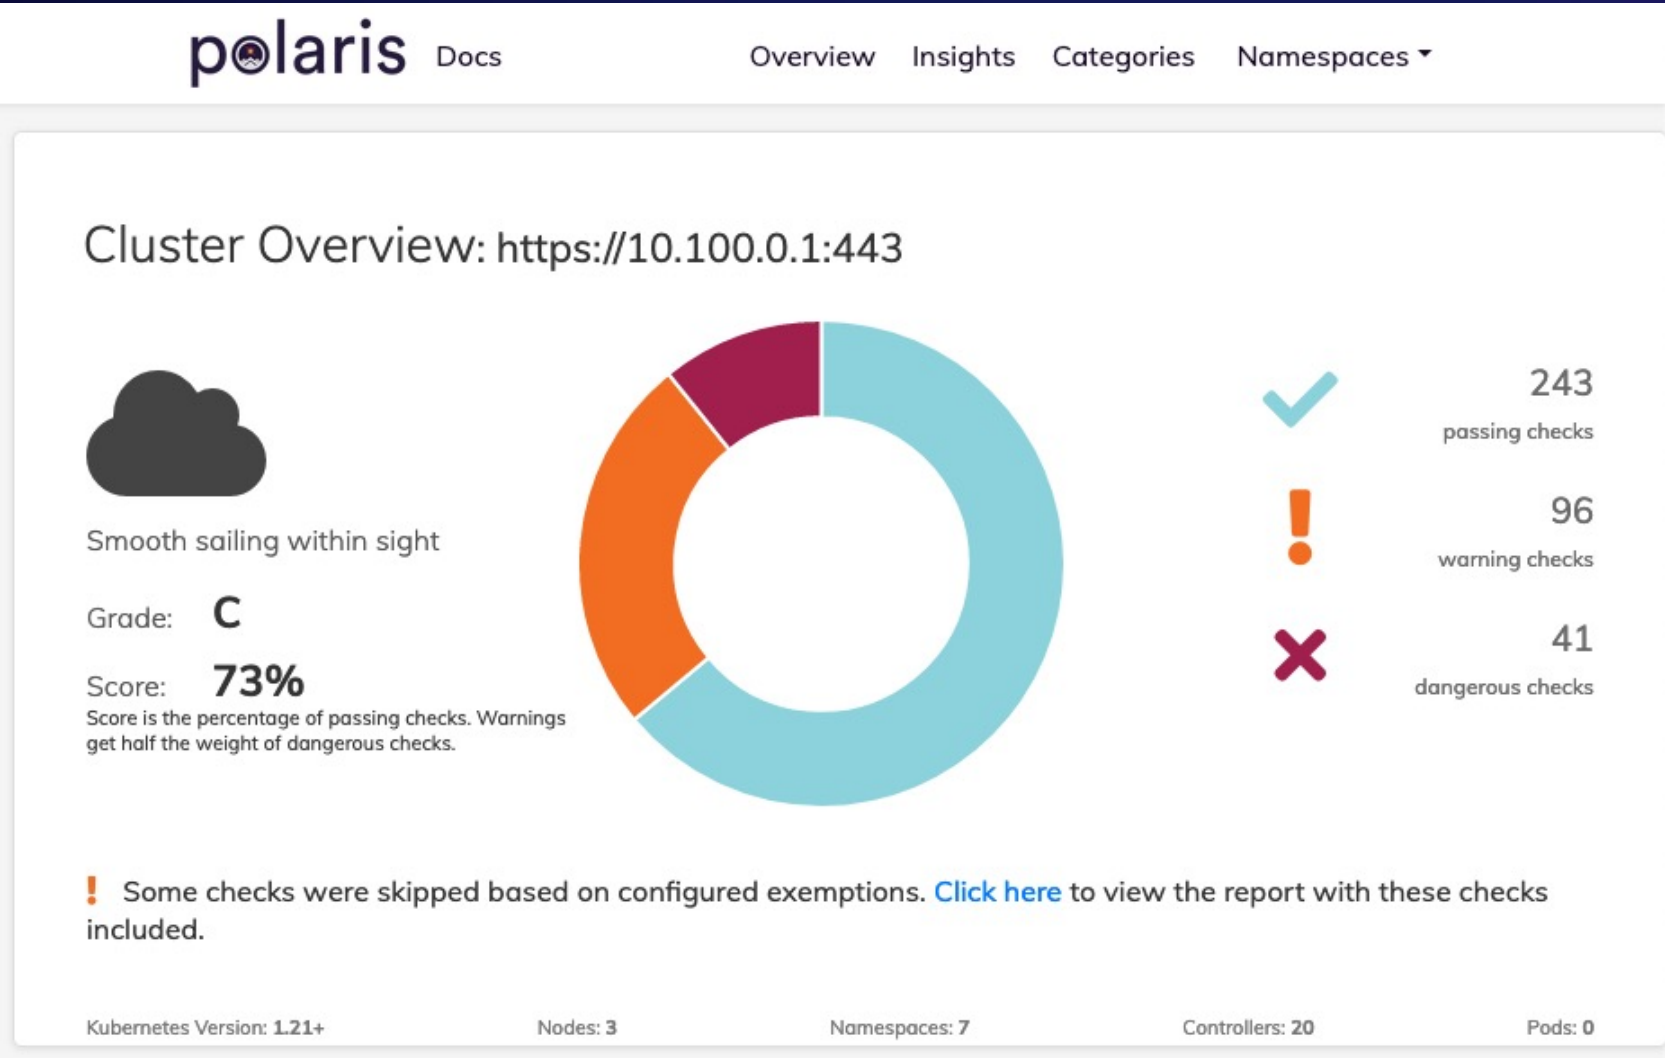
Pod and runtime security – Falco

```
[ec2-user@ip-10-0-1-211 wordpress]$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
falco-25rvh                         1/1     Running   2           4d10h
falco-4w2gb                         1/1     Running   1           4d10h
falco-nnk4f                         1/1     Running   1           4d10h
falco-rlf7j                         1/1     Running   2           4d10h
falco-wsv7d                         1/1     Running   1           4d10h
fluentbit-29sz8                     1/1     Running   7           4d10h
fluentbit-6nsw9                     1/1     Running   7           4d10h
fluentbit-9smnq                     1/1     Running   7           4d10h
fluentbit-lxx4z                     1/1     Running   1           4d10h
fluentbit-rk45t                     1/1     Running   7           4d10h
nginx2-78848c9dcb-5ssrm             1/1     Running   1           4d10h
nginx2-78848c9dcb-dsnch             1/1     Running   1           4d10h
nginx2-78848c9dcb-r7z64             1/1     Running   1           4d11h
```

```
[ec2-user@ip-10-0-1-211 wordpress]$ kubectl exec -it nginx2-78848c9dcb-5ssrm -- /bin/sh
# touch /etc/hacker
# cd /bin/ ; mkdir tools
# cat /etc/shadow > /dev/null 2>&1
```

▼	2021-07-02T10:39:34.401+09:00	{"log":"01:39:34.377084653: Notice A shell was spawned in a container with an attached terminal (user=root user_loginuid=-1 k8s.ns=default k8s.pod=nginx2-78848c9dcb-5ssrm container=b62164264cee shell=sh parent=runc cmdline=sh terminal=34...	Copy
▼	2021-07-02T10:39:45.128+09:00	{"log":"01:39:45.118927771: Error File below /etc opened for writing (user=root user_loginuid=-1 command=touch /etc/hacker parent=sh pcmdline=sh file=/etc/hacker program=touch gparent=\u003cNA\u003e ggparent=\u003cNA\u003e gggparent=\u003cNA\u003e container_id=b62164264cee image=nginx) k8s.ns=default k8s.pod=nginx2-78848c9dcb-5ssrm container=b62164264cee k8s.ns=default k8s.pod=nginx2-78848c9dcb-5ssrm container=b62164264cee", "stream": "stdout"}	Copy
▼	2021-07-02T10:41:03.340+09:00	{"log":"01:41:03.323755803: Error Directory below known binary directory created (user=root user_loginuid=-1 command=mkdir tools directory=/bin/tools container_id=b62164264cee image=nginx) k8s.ns=default k8s.pod=nginx2-78848c9dcb-5ssrm container=b62164264cee k8s.ns=default k8s.pod=nginx2-78848c9dcb-5ssrm container=b62164264cee", "stream": "stdout"}	Copy
▼	2021-07-02T10:41:32.865+09:00	{"log":"01:41:32.840168387: Warning Sensitive file opened for reading by non-trusted program (user=root user_loginuid=-1 program=cat command=cat /etc/shadow file=/etc/shadow parent=sh gparent=\u003cNA\u003e ggparent=\u003cNA\u003e gggparent=\u003cNA\u003e container_id=b62164264cee image=nginx) k8s.ns=default k8s.pod=nginx2-78848c9dcb-5ssrm container=b62164264cee k8s.ns=default k8s.pod=nginx2-78848c9dcb-5ssrm container=b62164264cee", "stream": "stdout"}	Copy

Pod and runtime security – Polaris by Fairwinds



Pod and runtime security

- Implement IRSA (IAM roles for service accounts)
- Run dynamic scans - Falco, Prisma Twistlock, Aqua, etc.
- Use Pod Security Standards or Policy as code
 - Deny privileged escalation
 - Deny running as root
 - Drop Linux capabilities
 - Implement OPA with Gatekeeper
 - Use other tools like Kyverno

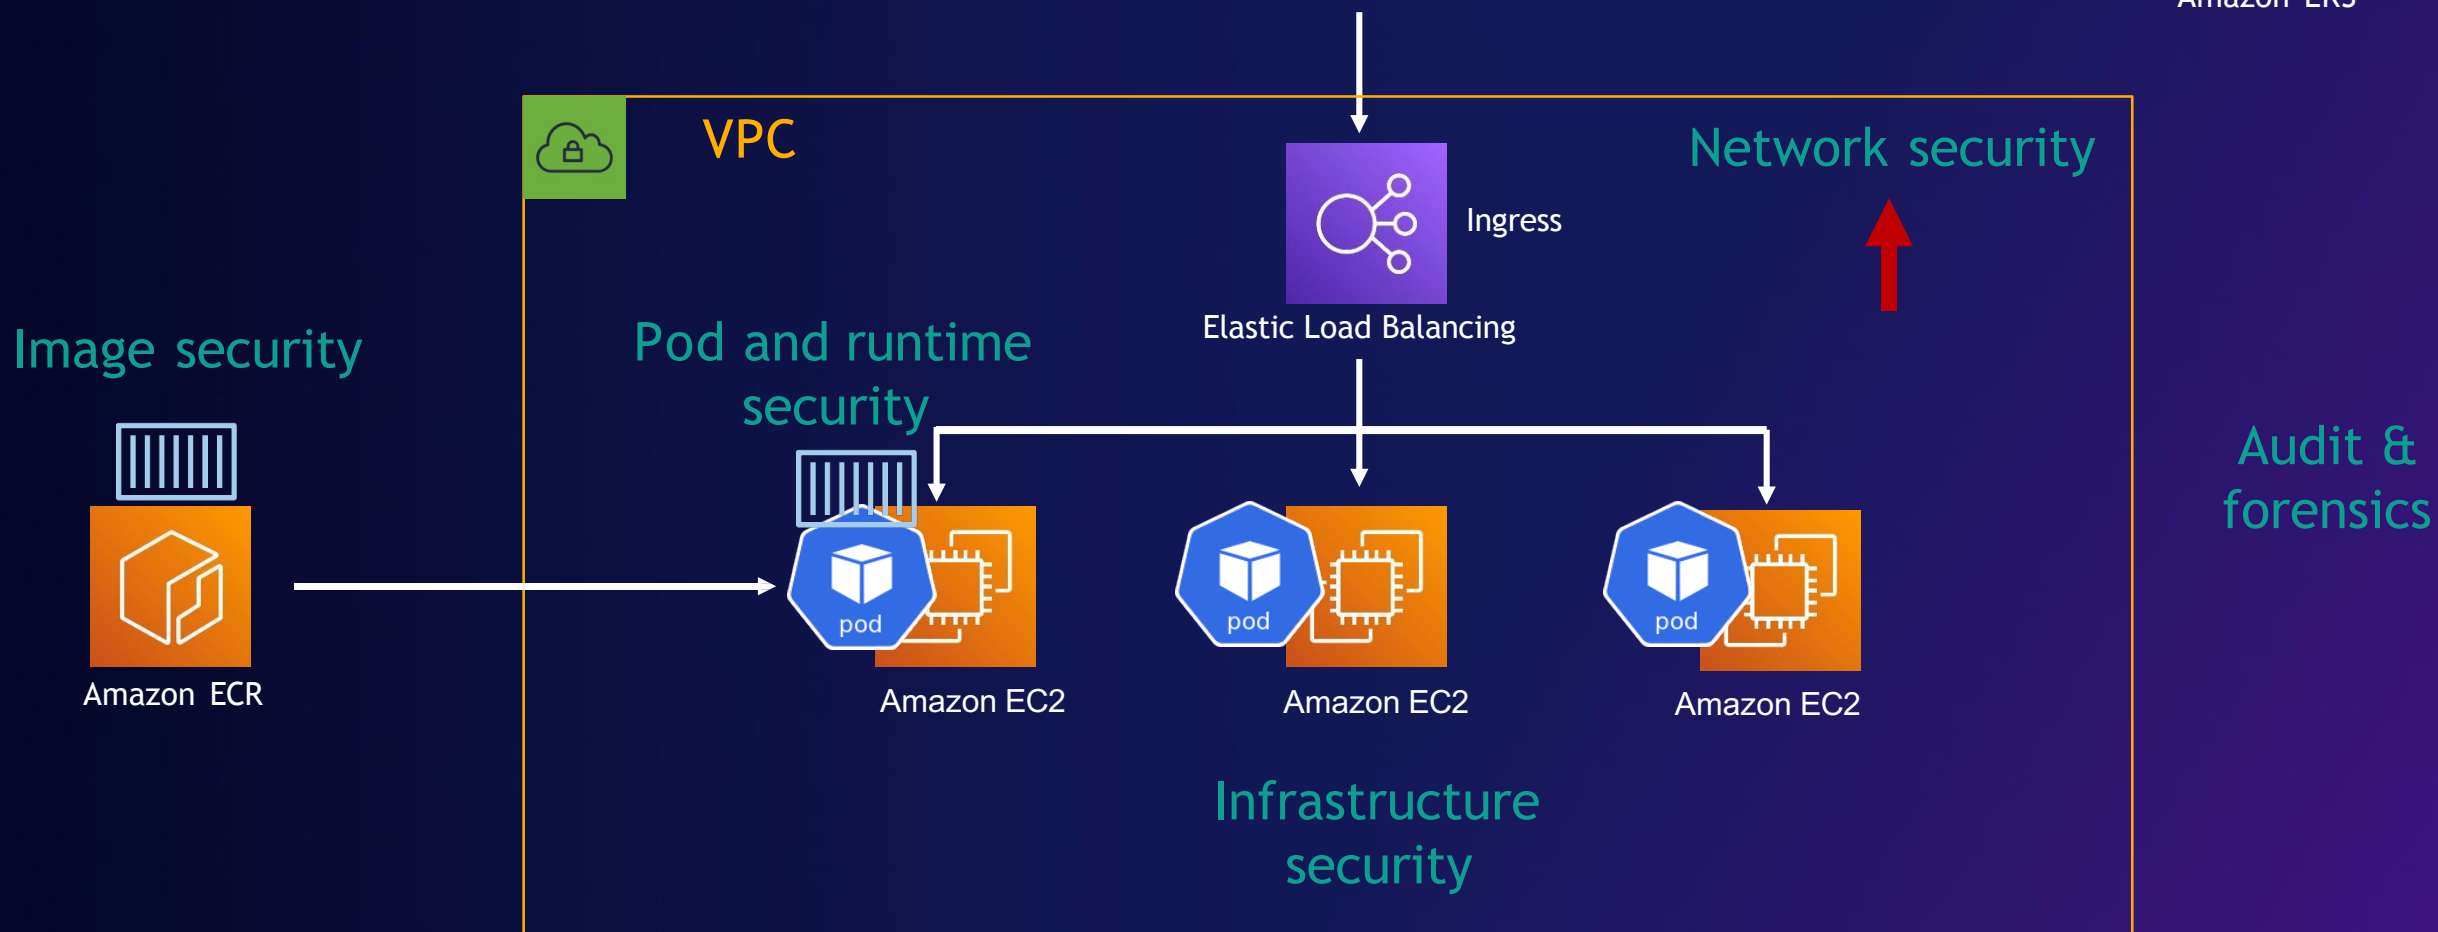
Pod and runtime security - Container and pods

- Dedicate service account for each application
- Do not allow privileged escalation
- Containers must be required to run as non-root
- Restrict the containers that can run as privileged
- Drop capabilities - default capabilities are not as strict as you think

```
apiVersion: v1
kind: Pod
metadata:
  name: security-pod
spec:
  serviceAccountName: sec-pod
  containers:
  - name: my-pod
    image: node
    securityContext:
      allowPrivilegeEscalation: false
      runAsUser: 1000
      runAsNonRoot: true
      privileged: false
      capabilities:
        drop:
        - all
```



Amazon EKS

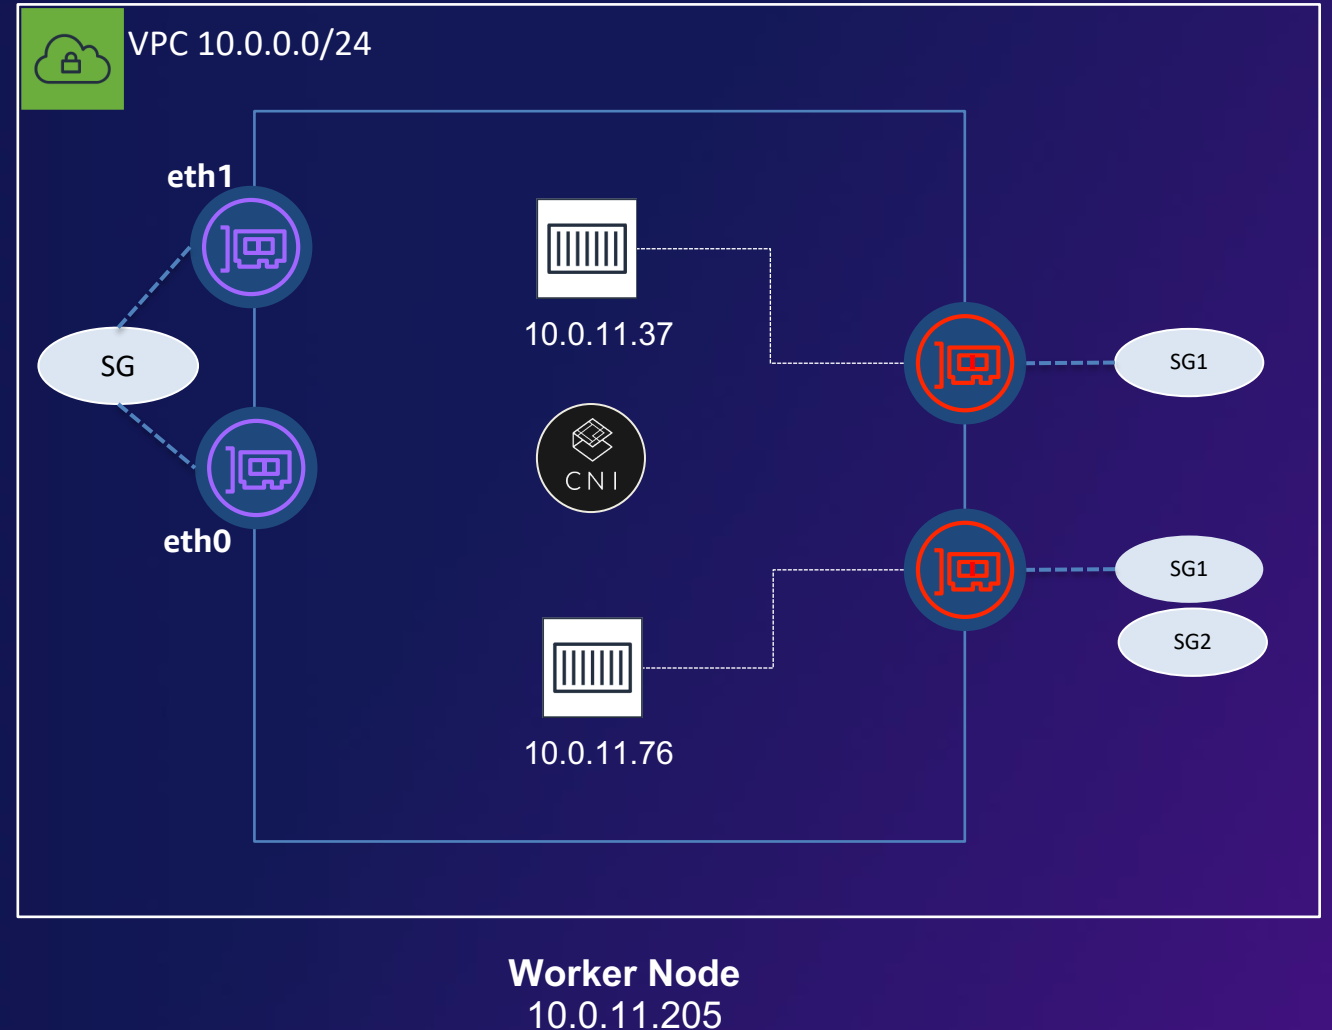


Network security

- Use security groups for pods to restrict traffic to AWS resources

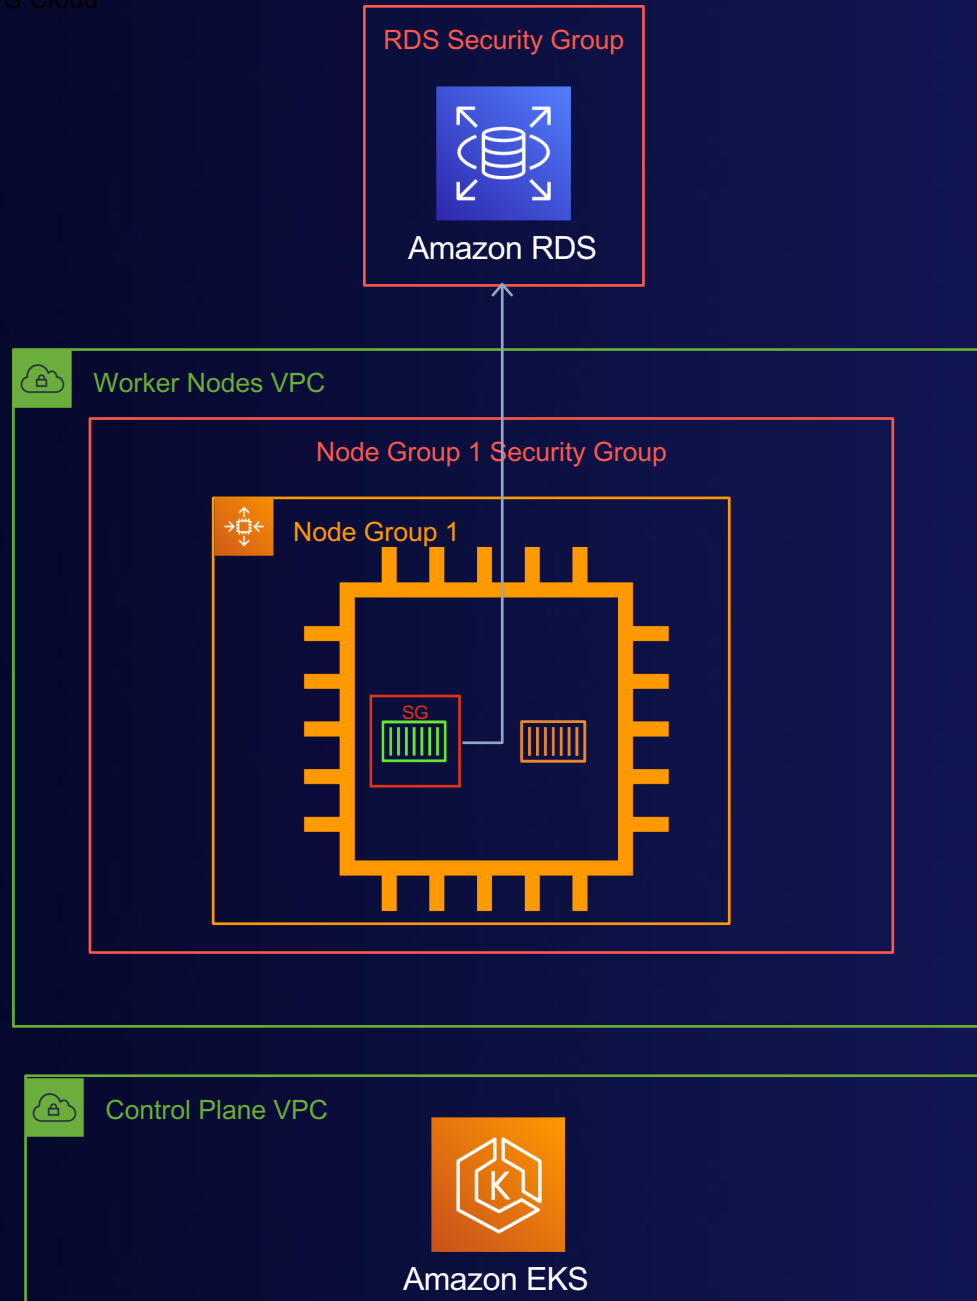
Securing Pods with EC2 Security Groups

- Worker node and all its Pods share the same security group(s) attached to the primary ENI
- Assign each Pod to a dedicated ENI with its own separate set of security groups(s)





AWS Cloud



RDS Security Group

Type	Protocol	Ports	Source
PostgreSQL	TCP	5432	sg-yyyyyyy (POD)

API Extension with new CRD

```
apiVersion: vpcresources.k8s.aws/v1beta1
kind: SecurityGroupPolicy
metadata:
  name: RDSClient
spec:
  serviceAccountSelector:
    matchLabels:
      role: RDSClient
  securityGroups:
    groupIds:
      - sg-yyyyyy
```



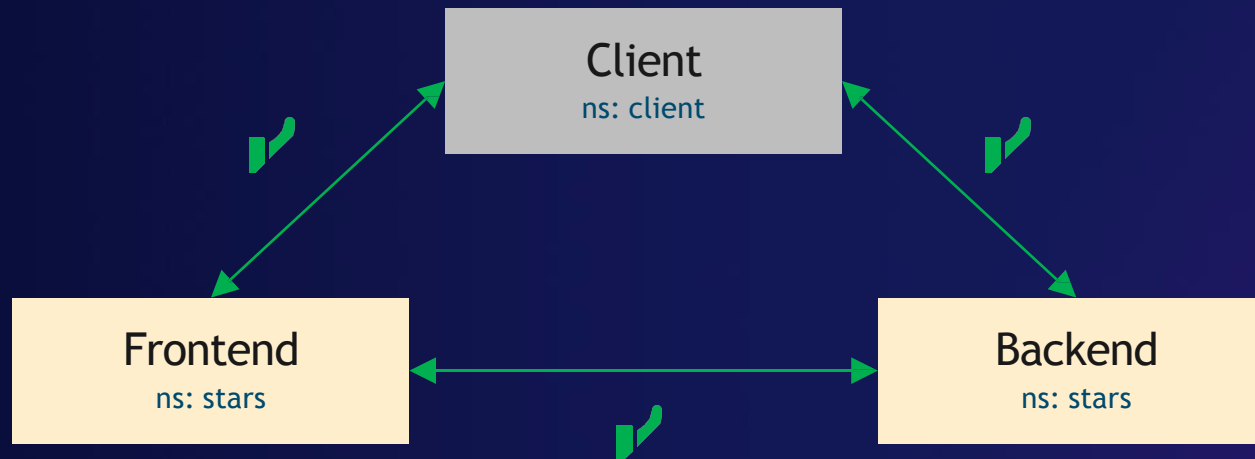
Amazon EKS

Network security

- Use security groups for pods to restrict traffic to AWS resources
- Use k8s network policies to restrict traffic within cluster

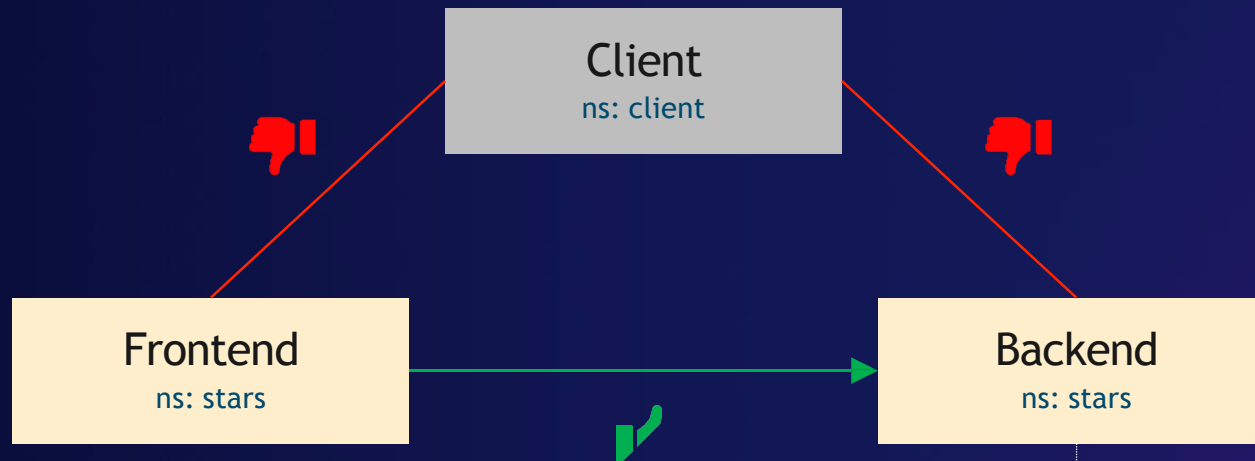
Network policy example

- Initial state without any network policy in effect



Network policy: Ingress rule

- Network policy to allow traffic to ingress into **Backend** pods from **Frontend** pods



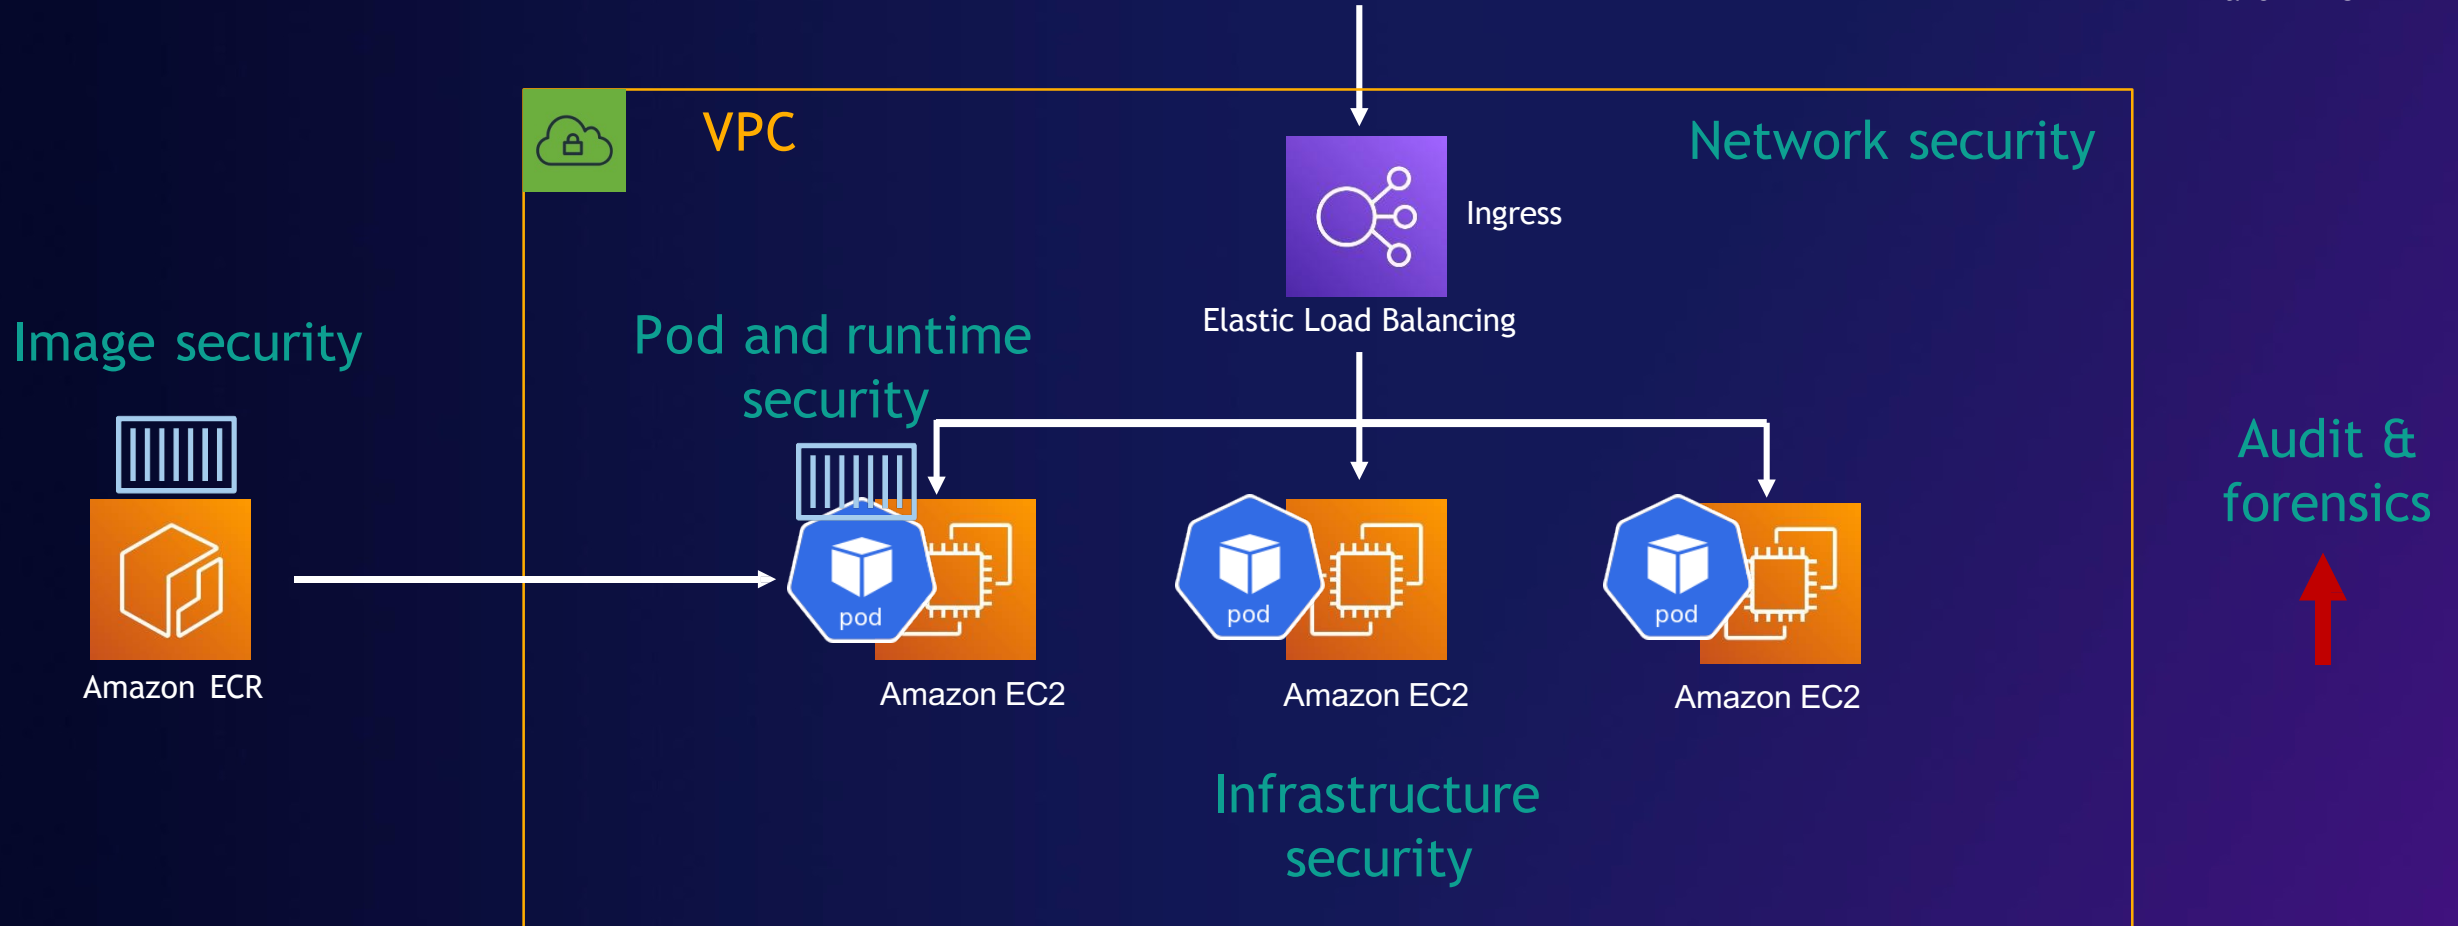
```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  namespace: stars
  name: backend-policy
spec:
  podSelector:
    matchLabels:
      role: backend
  ingress:
    - from:
        - podSelector:
            matchLabels:
              role: frontend
      ports:
        - protocol: TCP
          port: 6379
```

Network security

- Use security groups for pods to restrict traffic to AWS resources
- Use k8s network policies to restrict traffic within cluster
 - Enforce using network policy engine like Calico
- Encryption in transit
 - Terminate HTTPS external traffic at ELB
 - Use service mesh like App Mesh for mTLS service-service communication




Amazon EKS




Auditing and forensics

- Enable control plane logs
- Stream logs from containers to external log aggregator
- Periodically audit control plane and AWS CloudTrail logs for suspicious activity
- Immediately isolate pods you suspect have been compromised
 - Remove/change labels
 - Create network policy to isolate the pod
- Cordon the instance (if necessary)
 - Capture volatile artifacts on the worker node, e.g., memory, disk, etc.

EKS security best practices guide

 **EKS Best Practices Guides**

 Search

EKS Best Practices Guides
Guides >
Security ▾
[Home](#)
Identity and Access Management
Pod Security
Multi-tenancy
Detective Controls
Network Security
Data Encryption and Secrets Management
Runtime Security
Infrastructure Security
Regulatory Compliance
Incident Response and Forensics
Image Security

Amazon EKS Best Practices Guide for Security

This guide provides advice about protecting information, systems, and assets that are reliant on EKS while delivering business value through risk assessments and mitigation strategies. The guidance herein is part of a series of best practices guides that AWS is publishing to help customers implement EKS in accordance with best practices. Guides for Performance, Operational Excellence, Cost Optimization, and Reliability will be available in the coming months.

How to use this guide

This guide is meant for security practitioners who are responsible for implementing and monitoring the effectiveness of security controls for EKS clusters and the workloads they support. The guide is organized into different topic areas for easier consumption. Each topic starts with a brief overview, followed by a list of recommendations and best practices for securing your EKS clusters. The topics do not need to be read in a particular order.

<https://aws.github.io/aws-eks-best-practices/security/docs/>



Thank you!

